

Leveraged Learning

Daniel Chernowitz

August 19, 2026

This work rests on three insights.

1. Learning is the incremental ability to answer questions correctly that follow a pattern.
2. The number of questions reviewed is not the right measure of a learner's progress. What counts is the information received from asking them. Not all questions are created equal: about some, the learner already knows a great deal a priori, and their answers count for less.
3. One can infer *more* than a single bit from a yes–no question, by exploiting the correlations of an intelligent prior: an answer over here settles expectations over there. This ability to extrapolate is codified here as *leverage*.

Contents

1	A thought experiment	3
1.1	An unintelligent learner	3
1.2	Intelligent priors	5
1.2.1	A one-parameter world	6
2	The general case	8
2.1	The answer matrix and the update rule	9
2.2	Entropy bookkeeping	11
2.3	Playing the guiding example	12
2.4	Leverage	17
2.5	Correlators	17
2.5.1	General m	22
3	Expectations	24
3.1	Choosing the questions	27
3.2	The Gibbs complexity prior	28
3.2.1	NAND complexity, and two cheaper descriptors	28
3.2.2	The ensemble	30
3.2.3	Trajectories	31
4	The thermodynamic limit	33
4.1	A toy example: the spike prior and its renormalization flow	33
4.1.1	The two-state flow and exact leverage	34
4.1.2	The fixed- p thermodynamic limit	36
4.2	Thermodynamic-limit calculus	37
4.2.1	The finite sequence from which the limit is taken	37

4.2.2	A precise macroscopic-limit statement	38
4.3	A constructive complexity prior	39
4.3.1	Polynomials over $\mathbb{F}_2$, and the butterfly	40
4.3.2	Complexity classes and the sampling rule	41
4.3.3	Three terms for the leverage average	42
4.3.4	The macroscopic curve	43
4.3.5	Two statistics: $F(x) = x$ and $F(x) = x^2$	44
A	More characteristics of the spike prior	45
A.1	One question, branch by branch	45
A.2	The extraction ratio	46
A.3	The sharp-prior limit and the order of limits	47
B	Block priors and the Bernstein profile language	48
B.1	Independent blocks of questions	49
B.2	Block laws and their map priors	50
B.3	Bernstein polynomials as a profile language	51
B.4	Leverage curves at finite n	52

1 A thought experiment

Let us start with a small thought experiment, by introducing *Werner the learner* (Figure 1).




Figure 1: Werner the learner.

In this work, we model learning as the increased ability to give the correct answer to questions, as more of the ground truth is revealed. It is Werner’s job to learn a pattern. The most basic shape of a pattern is a *binary map*, and the most basic binary map is a map from one bit to one bit.

The cast of the experiment:

- a **question** $q \in \mathcal{Q} = \{0, 1\}$;
- an **answer** $a \in \mathcal{A} = \{0, 1\}$;
- the **truth** $\psi : \mathcal{Q} \rightarrow \mathcal{A}$, the one fixed map nature uses to answer questions. ψ is unknown to Werner. By definition then, $\psi(q) = a$.

There are exactly four candidate maps ϕ_j that ψ could be (Table 1): each of the two inputs can be sent to either of the two outputs, independently. We index them by reading the truth table $\phi_j(1)\phi_j(0)$ as a binary numeral: that number *is* j .

j	$\phi_j(1)$	$\phi_j(0)$	name
0	0	0	constant 0
1	0	1	NOT ($\neg q$)
2	1	0	identity (q)
3	1	1	constant 1

Table 1: The four binary maps from one bit to one bit.

1.1 An unintelligent learner

Now, if Werner is unintelligent, he has no preconceived notion of which maps are more useful, more likely, or somehow more logical than others. In practice, we would have to tell Werner the answer a to every single question q , before he has learned the whole pattern.

We model this with a **uniform prior**: writing $p_j = \Pr(\psi = \phi_j)$ for Werner’s belief, every candidate map carries the same weight.

EXAMPLE

For the four maps at hand:

$$p_j = \frac{1}{4}, \quad j = 0, 1, 2, 3. \quad (1)$$

What does this belief predict? For each question q separately, we can tabulate the probability of each possible answer a : it is the total belief in the maps that would answer a to q ,

$$M_{a,q} = \sum_{j: \phi_j(q)=a} p_j. \quad (2)$$

Arranged with one row per answer ($a = 0, 1$ from the top) and one column per question ($q = 1, 0$ from the left, descending like the digits of the numeral convention), these numbers form the matrix M - each column a probability distribution over answers, so M is column-stochastic.

EXAMPLE

Under the uniform prior, two of the four maps answer 0 and two answer 1, at either question:

$$M = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & M_{0,1} & M_{0,0} \\ \mathbf{1} & M_{1,1} & M_{1,0} \end{array} = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & p_0 + p_1 & p_0 + p_2 \\ \mathbf{1} & p_2 + p_3 & p_1 + p_3 \end{array} = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & \frac{1}{2} & \frac{1}{2} \\ \mathbf{1} & \frac{1}{2} & \frac{1}{2} \end{array}. \quad (3)$$

Every column is a fair coin: the unintelligent Werner predicts nothing about any answer before it is revealed.

Next, Werner gets to predict the answer to the question $q = 0$. By his own belief it is a coin toss: $M_{0,0} = M_{1,0} = \frac{1}{2}$. Let us say the answer comes back $a = 1$. The **surprisal** of this answer, to Werner, is 1 bit, because the probability he assigned to it was one half.

In general, when an event a belief holds with probability P actually occurs, its surprisal is

$$s = -\log_2 P \quad \text{bits}, \quad (4)$$

here $s = -\log_2 M_{a,q}$ for answer a arriving at question q . Surprisal measures how much the received information deviates from what was already expected: an answer held certain ($P = 1$) carries 0 bits — nothing new; a coin toss carries exactly 1 bit; and the surprisal grows without bound as the answer gets more unexpected, $P \rightarrow 0$. It is additive: the surprisal of two independent events occurring is the sum of their surprisals, just as their probabilities multiply.

Moreover, Werner can use a process of elimination — an extreme form of Bayes' rule — to update his prior into a posterior. Semantically: every candidate map that would have answered the question wrongly is now *ruled out*, not merely disfavored.

EXAMPLE

The truth ψ answered 1 at $q = 0$, and a map with $\phi_j(0) = 0$ cannot be ψ . This kills the constant-0 map ($j = 0$) and the identity ($j = 2$). The two survivors, NOT ($j = 1$) and constant 1 ($j = 3$), both predicted the observed answer, so this round of evidence does not distinguish between them at all (Table 2).

j	$\phi_j(1)$	$\phi_j(0)$	name
0	0	0	constant 0
1	0	1	NOT ($\neg q$)
2	1	0	identity (q)
3	1	1	constant 1

Table 2: The candidate maps after the answer $\psi(0) = 1$: the two maps with $\phi_j(0) = 0$ are ruled out.

The surviving ratio of believabilities must be preserved, and the only update that does so is to renormalize the surviving weights to sum to one. (This is Bayes' rule with a 0/1 likelihood: $\Pr(\psi = \phi_j \mid a) \propto \Pr(a \mid \phi_j) p_j$, and $P(a \mid \phi_j)$ is 1 for a map that answers a and 0 for a map that does not.)

EXAMPLE

The posterior is

$$p' = (0, \frac{1}{2}, 0, \frac{1}{2}). \quad (5)$$

Constructing the new answer matrix from p' by the same recipe as before:

$$M' = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & p'_0 + p'_1 & p'_0 + p'_2 \\ \mathbf{1} & p'_2 + p'_3 & p'_1 + p'_3 \end{array} = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & \frac{1}{2} & 0 \\ \mathbf{1} & \frac{1}{2} & 1 \end{array}. \quad (6)$$

So indeed, Werner knows the answer to $q = 0$ — that column has collapsed onto $a = 1$ — and he is still as agnostic as possible about the other answer: the $q = 1$ column remains a fair coin. He gained 1 bit of surprisal information, and cleaned up his table to the tune of 1 bit: the $q = 0$ column went from a full bit of answer-uncertainty to none, while the $q = 1$ column is unchanged.

After the next question, $q = 1$, he learns $a = 1$ — again a coin toss beforehand ($M'_{1,1} = \frac{1}{2}$), so again worth 1 bit of surprisal. Elimination now rules out NOT ($\phi_1(1) = 0$), leaving a single survivor:

$$p'' = (0, 0, 0, 1), \quad M'' = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & 0 & 0 \\ \mathbf{1} & 1 & 1 \end{array}. \quad (7)$$

So this is the constant map: $\psi = \phi_3$. Werner is done. Every column is one-hot, and he will be able to answer all questions correctly from now on.

1.2 Intelligent priors

Now, let us assume that Werner knows *something* about the world before he starts.

EXAMPLE

Let us say Werner knows the map is not constant. His belief then spreads evenly over the two remaining candidates, NOT ($j = 1$) and the identity ($j = 2$):

$$p = (0, \frac{1}{2}, \frac{1}{2}, 0). \quad (8)$$

Is he any better at predicting $\psi(0)$? Constructing M by the usual recipe:

$$M = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & p_0 + p_1 & p_0 + p_2 \\ \mathbf{1} & p_2 + p_3 & p_1 + p_3 \end{array} = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & \frac{1}{2} & \frac{1}{2} \\ \mathbf{1} & \frac{1}{2} & \frac{1}{2} \end{array}. \quad (9)$$

He is still just as bad as before: every column a fair coin.

So what has all this veteran wisdom about the nature of patterns in the world bought him? His increased ability for *inference*.

EXAMPLE

Ask $q = 0$ again, and let the answer again come back $a = 1$ — again a coin toss beforehand, so again worth exactly 1 bit of surprisal. Elimination rules out the identity ($\phi_2(0) = 0$), and the lone survivor is NOT:

$$p' = (0, 1, 0, 0), \quad M' = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & 1 & 0 \\ \mathbf{1} & 0 & 1 \end{array}. \quad (10)$$

This time *both* columns collapse at once. Werner received the same single bit of surprisal as before, but cleaned up 2 bits of table: the answer at $q = 0$ he observed, and the answer at $q = 1$ he *deduced*, without ever asking — ruling out the constants means $\phi(1) \neq \phi(0)$, so his prior held the two answers perfectly anticorrelated, and settling one settles the other. He is done after one question instead of two.

We say the information from the first answer was **leveraged**. We will make this more precise in the upcoming sections.

1.2.1 A one-parameter world

For more structure, let us relax the claim that there are *no* constant maps. Instead, we create a one-parameter world where the prior probability that the map is constant is p , spread evenly within each class (Table 3):

j	$\phi_j(1)$	$\phi_j(0)$	name	p_j
0	0	0	constant 0	$p/2$
1	0	1	NOT ($\neg q$)	$(1-p)/2$
2	1	0	identity (q)	$(1-p)/2$
3	1	1	constant 1	$p/2$

Table 3: The one-parameter prior: constant with probability p .

In vector form, and with the answer matrix constructed by the usual recipe:

$$(p_0, p_1, p_2, p_3) = \frac{1}{2}(p, 1-p, 1-p, p), \quad M = \begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & \frac{1}{2} & \frac{1}{2} \\ \mathbf{1} & \frac{1}{2} & \frac{1}{2} \end{array}, \quad (11)$$

fair coins in every column, for every value of p .

How much is there to learn, in total? We account for the table as a whole by declaring entropy *extensive* over the columns — the answer-uncertainties of the columns simply add:

$$H(M) := \sum_q H(M_{\cdot,q}), \quad H(M_{\cdot,q}) := - \sum_a M_{a,q} \log_2 M_{a,q}. \quad (12)$$

Here both columns are fair coins carrying 1 bit each, so the table starts at its maximum: $H(M) = 2$ bits.

Now let us ask a question. Again we could use $q = 0$; by the symmetry of the prior it actually does not matter for this argument. And we are now going to talk about *expected* values, instead of realized ones: averages over the possible maps, denoted $\langle \cdot \rangle$. Since we assume our prior is *a priori* correct — nature draws the truth from the very distribution Werner believes — the average over ψ simply zips with the weights: $\langle X \rangle = \sum_j p_j X(\psi=\phi_j)$.

How much surprisal should Werner expect from asking a question q ? The answer comes back a exactly when the truth answers a , which under the zipped average happens with probability $\sum_{j:\phi_j(q)=a} p_j = M_{a,q}$ — the very probability Werner assigned to it — and when it does, it costs him $-\log_2 M_{a,q}$ of surprisal. Grouping the average over maps by the answer each map gives:

$$\langle s_q \rangle = \sum_j p_j (-\log_2 M_{\phi_j(q),q}) = \sum_a \underbrace{\sum_{j:\phi_j(q)=a} p_j}_{M_{a,q}} (-\log_2 M_{a,q}) = \sum_a M_{a,q} (-\log_2 M_{a,q}). \quad (13)$$

The right-hand side is exactly the column entropy $H(M_{\cdot,q})$ of equation (12): **the expected surprisal of a question is the entropy of its column**. Each answer occurs with exactly the probability Werner assigned to it, so his average surprisal is his own uncertainty about that answer — an identity that will do a great deal of work in what follows.

For a two-answer column, this entropy is a function of a single number, the probability x of one of the two answers: the **binary entropy function**

$$h_2(x) := -x \log_2 x - (1-x) \log_2 (1-x), \quad (14)$$

zero at $x \in \{0, 1\}$ (a settled answer) and maximal, 1 bit, at the fair coin $x = \frac{1}{2}$. For the first question here the column is a fair coin, so

$$\langle s_1 \rangle = H(M_{\cdot,0}) = h_2(\tfrac{1}{2}) = 1 \text{ bit}, \quad (15)$$

whatever the value of p .

And the other column is now cleaned up. Breaking the update down over the possible maps: with probability $\frac{1}{2}$ the answer is 1 (truth is NOT or constant 1), with probability $\frac{1}{2}$ it is 0 (truth is FALSE or the identity), and

$$M' = \underbrace{\begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & 1-p & 0 \\ \mathbf{1} & p & 1 \end{array}}_{a_1=1, \text{ prob } 1/2} \quad \text{or} \quad \underbrace{\begin{array}{c|cc} \mathbf{a} \backslash \mathbf{q} & \mathbf{1} & \mathbf{0} \\ \hline \mathbf{0} & p & 1 \\ \mathbf{1} & 1-p & 0 \end{array}}_{a_1=0, \text{ prob } 1/2}. \quad (16)$$

Either way, the asked column is one-hot and the unasked column is a coin of bias p : what survives on $q = 1$ is precisely the class uncertainty — constant or not — and nothing else. The total entropy of the table is now

$$H(M') = 0 + h_2(p) = h_2(p), \quad (17)$$

a reduction of $2 - h_2(p)$ bits, purchased with a single bit of surprisal. Define the **leverage after one question** as the ratio of table entropy destroyed to surprisal received:

$$L_1 := \frac{2 - h_2(p)}{1} = 2 - h_2(p) \in [1, 2]. \quad (18)$$

The expected surprisal on the second question is, once more, the entropy of its column,

$$\langle s_2 \rangle = h_2(p) \leq 1 \text{ bit}, \quad (19)$$

after which the table is empty. Cumulatively over the two questions, $1 + h_2(p)$ bits of received information reduced the table by the full 2 bits it started with, so the **cumulative leverage** after two questions is

$$L_2 := \frac{2}{1 + h_2(p)} \in [1, 2]. \quad (20)$$

So our knowledge about the world has allowed us to be more certain after our inference. Set $p = \frac{1}{2}$, and the prior is the uniform one of Section 1.1: $h = 1$, $L_1 = L_2 = 1$, and nothing is leveraged at all. As p skews further from $\frac{1}{2}$, the second question carries less and less expected information, and Werner’s potential for learning from it goes down with it. The column it could clean up is nearly settled already. In the limit $p = 0$, we recover the non-constant world from the start of Section 1.2. There is no surprisal and no improvement to M left in the second question at all: it contributes nothing to denominator or numerator, and $L_1 = L_2$. Figure 2 traces all three curves between these extremes.




Figure 2: The one-parameter world as a function of the constant-map probability p . Left: the table entropy $h_2(p)$ remaining after one question, in bits. Right: the one-question leverage L_1 and the cumulative leverage L_2 after both questions. At $p = \frac{1}{2}$ (the uniform prior) nothing is leveraged; toward the deterministic-class extremes every answer doubles its value.

In general, in a larger world, if there are d questions to disagree on between the two maps, the reader can convince themselves that the limit as $p \rightarrow 1$ is $L_1 = d$.

2 The general case

Now let us give the definitions in full generality. Questions are bit strings of length n , and answers are bit strings of length m :

$$q \in \mathcal{Q} = \{0, 1\}^n, \quad a \in \mathcal{A} = \{0, 1\}^m, \quad (21)$$

and we write the sizes of those two sets without bars,

$$Q := |\mathcal{Q}| = 2^n, \quad A := |\mathcal{A}| = 2^m, \quad (22)$$

so there are Q distinct questions and A possible answers to each. As before, we read bit strings as the numbers they spell, using the two interchangeably: $q \in \{0, \dots, Q-1\}$ and $a \in \{0, \dots, A-1\}$. The truth $\psi : \mathcal{Q} \rightarrow \mathcal{A}$ is one of the candidate maps ϕ_j ; since each of the Q inputs can be sent to any of the A outputs independently, there are

$$N = A^Q = 2^{mQ} \quad (23)$$

of them, and the index j is again the truth table read as a numeral, now with Q digits in base A : $j = \sum_q \phi_j(q) A^q$, running from the constant-0 map ($j = 0$) to the constant- $(A-1)$ map ($j = N-1$). Werner's belief is a distribution over all of them, $\sum_{j=0}^{N-1} p_j = 1$.

EXAMPLE

The example we will use to guide the general case is $n = 2$, $m = 1$, which has $N = 2^{1 \cdot 4} = 16$ maps: every Boolean function of the two input bits b_1, b_0 , with $q = b_1 b_0$ read as an ordinary binary number, $q = 2b_1 + b_0$. Each ϕ_j in Table 4 is shown as its truth table, printed in descending question order (11, 10, 01, 00) so that the answer to question q occupies the digit of weight 2^q . The printed string is exactly j in binary (e.g. AND, which answers 1 only to $q = 11$, reads 1000, so it is $j = 8$). The prior p_j alongside is conjured for illustration.

j	$\phi_j(q)$				name	p_j	
	11	10	01	00			
0	0	0	0	0	FALSE (constant 0)	0.394	
1	0	0	0	1	$\neg(b_1 \vee b_0)$ (NOR)	0.010	
2	0	0	1	0	$\neg b_1 \wedge b_0$	0.010	
3	0	0	1	1	$\neg b_1$	0.010	
4	0	1	0	0	$b_1 \wedge \neg b_0$	0.131	
5	0	1	0	1	$\neg b_0$	0.010	
6	0	1	1	0	$b_1 \oplus b_0$ (XOR)	0.010	
7	0	1	1	1	$\neg(b_1 \wedge b_0)$ (NAND)	0.010	
8	1	0	0	0	$b_1 \wedge b_0$ (AND)	0.213	
9	1	0	0	1	$b_1 = b_0$ (XNOR)	0.092	
10	1	0	1	0	b_0 (projection)	0.040	
11	1	0	1	1	$b_1 \rightarrow b_0$ (IF THEN)	0.010	
12	1	1	0	0	b_1 (projection)	0.010	
13	1	1	0	1	$b_0 \rightarrow b_1$ (IF THEN)	0.030	
14	1	1	1	0	$b_1 \vee b_0$ (OR)	0.010	
15	1	1	1	1	TRUE (constant 1)	0.010	

Table 4: The guiding example: all 16 maps of the $n = 2$, $m = 1$ hypothesis space, with a contrived prior ($\sum_j p_j = 1.000$).

2.1 The answer matrix and the update rule

The answer matrix is defined by the same rule as before — the total belief in the maps that would answer a to q :

$$M_{a,q} = \sum_{j: \phi_j(q)=a} p_j, \quad (24)$$

now a $A \times Q$ matrix, still with each column a probability distribution over the possible answers: column-stochastic.

It is worth taking a moment to describe the constitution of this matrix. Because of the enumeration of j (the index *is* the truth table) the answer matrix has a natural structure: it is the contraction of the belief vector p with a fixed tensor with $\{0,1\}$ entries E , summed along the map index,

$$M_{a,q} = \sum_{j=0}^{N-1} E_{a,q,j} p_j, \quad E_{a,q,j} := \delta(\phi_j(q), a) = \delta([j \cdot 2^{-qm}] \bmod A, a). \quad (25)$$

The second form is pure arithmetic on the index: by the numeral convention, $\phi_j(q)$ is nothing but the q -th digit of j in base A , so E never needs to be stored. Membership of j in the entry $M_{a,q}$ is read directly off the digits of j . This is a good thing, E is astronomically large for even moderate m and n . For $m = 1$ this reduces to: bit q of j equals a . The tensor E is moreover sparse and perfectly regular: each question-slice (q fixed) holds exactly N nonzero entries: every map answers exactly one a divided evenly over the A answer rows, N/A apiece. Within a slice the nonzero row follows the odometer pattern of a numeral system: runs of A^q consecutive j share a digit, cycling through all A rows with period A^{q+1} .

The odometer pattern gives every entry of M the same compact closed form: for $m = 1$, bit q of j equals 0 exactly when $j \bmod 2^{q+1} < 2^q$, and equals 1 otherwise.

EXAMPLE

Each entry of M collects the eight maps of Table 4 whose truth table carries the right digit. Filling the entries in — note the period halving column by column, with the outer columns written in their simplified forms, $j \bmod 16 < 8$ being simply $j < 8$ and $j \bmod 2 < 1$ simply “ j even” — and then substituting the values of p_j :

$$M = \begin{array}{c|cccc} \mathbf{a \backslash q} & \mathbf{11} & \mathbf{10} & \mathbf{01} & \mathbf{00} \\ \hline \mathbf{0} & \sum_{j < 8} p_j & \sum_{j \bmod 8 < 4} p_j & \sum_{j \bmod 4 < 2} p_j & \sum_{j \text{ even}} p_j \\ \mathbf{1} & \sum_{j \geq 8} p_j & \sum_{j \bmod 8 \geq 4} p_j & \sum_{j \bmod 4 \geq 2} p_j & \sum_{j \text{ odd}} p_j \end{array} \quad (26)$$

$$= \begin{array}{c|cccc} \mathbf{a \backslash q} & \mathbf{11} & \mathbf{10} & \mathbf{01} & \mathbf{00} \\ \hline \mathbf{0} & 0.585 & 0.779 & 0.890 & 0.818 \\ \mathbf{1} & 0.415 & 0.221 & 0.110 & 0.182 \end{array} .$$

The update rule also carries over verbatim. When question q is asked and the answer $a = \psi(q)$ comes back, the process of elimination reads, in one formula,

$$p'_j = \frac{\delta(\phi_j(q), a) p_j}{M_{a,q}} = \frac{E_{a,q,j} p_j}{M_{a,q}}. \quad (27)$$

The Kronecker delta performs the elimination. Only maps that agree with a survive and the division renormalizes the survivors, whose total prior mass is exactly $M_{a,q}$ by equation (24). Note that one and the same number prices the answer and funds the update: the event arrives with the probability $M_{a,q}$, costs Werner $-\log_2 M_{a,q}$ bits of surprisal, and its belief mass becomes the new unit of normalization. Every update shrinks the support by A^{-1} , radically thinning, especially for large m . The survivors live on a block-wave in j -space. The map determines the offset and frequency. The total set of survivors depends on the set of asked questions, not the order: updating commutes. This is illustrated in the example below.

EXAMPLE

Figure 3 makes the process visual for $n = m = 2$: a schematic prior over $N = 4^4 = 256$ maps, drawn as vertical lines with thickness representing weight. Each answer stops 3 out of every 4 surviving lines: maps whose truth table carry a wrong digit at the asked question. Renormalization inflates the weight of the rest, so the total probability is always unit. Four questions cut $256 \rightarrow 64 \rightarrow 16 \rightarrow 4 \rightarrow 1$, and the last line standing is the truth. Note the regularity of the cull: each answer keeps a single residue class of j , so the cadence of the survivors is regular.




Figure 3: The update rule as a cull, for $n = m = 2$. All 256 candidate maps start as vertical lines under a schematic prior. The truth is the bitwise negation $\psi(q) = \neg q$ ($j = 27$), and the four questions are asked in the order 00, 10, 01, 11. Each answer terminates the 3/4 of surviving maps whose truth-table digit at the asked question disagrees. Every survivor's thickness, proportional to its weight, is rescaled by the same renormalization factor, preserving their ratios. After four questions a single map remains. Interestingly, the first cull's survivors are evenly spaced because it's the smallest question. After the three lowest q , even spacing returns.

2.2 Entropy bookkeeping

The total entropy of the table is extensive over the columns, exactly as in equation (12):

$$H(M) = \sum_{q \in Q} H(M_{\cdot, q}) = - \sum_{q \in Q} \sum_{a \in A} M_{a, q} \log_2 M_{a, q} \leq mQ, \quad (28)$$

each column carrying at most m bits (a uniform spread over its A answers). The ceiling mQ is the size of the entire truth table in bits, and it is attained precisely by the uniform prior: maximal ignorance prices every answer at full cost.

The belief itself carries an entropy as well, and we give its gap to the table a name:

$$H(p) = - \sum_{j=0}^{N-1} p_j \log_2 p_j, \quad C := H(M) - H(p) \geq 0. \quad (29)$$

$H(p)$, Werner’s uncertainty about the *identity* of the truth, obeys the same ceiling: a distribution over N maps has entropy at most $\log_2 N = mQ$, attained again by the uniform prior. The information in the prior is not directly useful for prediction, but it is good to track.

Why is $C > 0$? The answers jointly determine the map, so $H(p)$ is the *joint* entropy of all Q answers, while $H(M)$ sums their *marginal* entropies and a joint entropy never exceeds the sum of its marginals. Hence $H(p) \leq H(M)$, and the gap C , the **total correlation** of the answers, is nonnegative: knowledge stored in correlations *between* answers, invisible to every single column of M .

We choose this convention for $H(M)$ because in the ignorant case

1. $H(M)$ is the total amount of information we must collect from questions to settle the map
2. $H(M) = H(p)$, so the ignorance agrees in both bases, without correlations.

Therefore, the scaling is sensible. How these quantities respond to an intelligent prior, is the main question of this work.

2.3 Playing the guiding example

Let us watch these ledgers in motion by playing the guiding example through all four questions. Werner can compute two forecasts for every unasked question, in advance, from M alone: the expected surprisal $\langle s \rangle$, which by equation (13) is the entropy of that question’s column, and the expected drop $\langle \Delta H(M) \rangle$ of the *whole* table if that question is asked next. This is the average over possible a , weighted by their probability (according to M):

$$\langle s \rangle(q) = \sum_a M_{a,q} (-\log_2 M_{a,q}), \quad (30)$$

$$\langle \Delta H(M) \rangle(q) = \sum_a M_{a,q} [H(M) - H(M^{(q,a)})], \quad (31)$$


where $M^{(q,a)}$ is the table after the update rule (27) processes answer a to question q . We know the column of q will drop $\langle s \rangle(q)$ in entropy. But other columns will also respond. The excess of the second over the first is the expected transfer to unasked columns, and it can only be paid out of correlations. We append both forecasts as extra rows under M .

EXAMPLE

Werner plays greedily: always asking the question with the largest forecast. The truth is AND ($\psi = \phi_8$).

Question 1. The starting state:

$\mathbf{a \backslash q}$	11	10	01	00
0	0.585	0.779	0.890	0.818
1	0.415	0.221	0.110	0.182
$\langle s \rangle$	0.979	0.762	0.500	0.684
$\langle \Delta H(M) \rangle$	1.104	0.800	0.544	0.801



To see where the forecast entries come from, take $q = 11$: the answer $a = 0$ arrives with probability $M_{0,11} = 0.585$ and would drop $H(M)$ from 2.925 to 1.569; the answer $a = 1$ arrives with probability 0.415 and drops it only to 2.178, so

$$\begin{aligned}\langle s \rangle(11) &= 0.585 (0.773) + 0.415 (1.269) = 0.979, \\ \langle \Delta H(M) \rangle(11) &= 0.585 (1.357) + 0.415 (0.748) = 1.104.\end{aligned}$$

Both forecast rows point to $q = 11$: it is the closest thing to a coin on offer ($\langle s \rangle = 0.979$) and also moves the whole table most, by a wide margin. Werner asks it, and nature answers $a = 1$ (probability 0.415, surprisal $s_1 = 1.269$). Eight maps die:

j	$\phi_j(11)$	name	p_j	p'_j	
0	0	FALSE (constant 0)	0.394		
1	0	$\neg(b_1 \vee b_0)$ (NOR)	0.010		
2	0	$\neg b_1 \wedge b_0$	0.010		
3	0	$\neg b_1$	0.010		
4	0	$b_1 \wedge \neg b_0$	0.131		
5	0	$\neg b_0$	0.010		
6	0	$b_1 \oplus b_0$ (XOR)	0.010		
7	0	$\neg(b_1 \wedge b_0)$ (NAND)	0.010		
8	1	$b_1 \wedge b_0$ (AND)	0.213	0.513	
9	1	$b_1 = b_0$ (XNOR)	0.092	0.222	
10	1	b_0 (projection)	0.040	0.096	
11	1	$b_1 \rightarrow b_0$ (IF THEN)	0.010	0.024	
12	1	b_1 (projection)	0.010	0.024	
13	1	$b_0 \rightarrow b_1$ (IF THEN)	0.030	0.072	
14	1	$b_1 \vee b_0$ (OR)	0.010	0.024	
15	1	TRUE (constant 1)	0.010	0.024	

The ledger:

$$H(M): 2.925 \rightarrow 2.178, \quad H(p): 2.707 \rightarrow 2.093, \quad C: 0.218 \rightarrow 0.085.$$


Werner received 1.269 bits but destroyed only 0.748 bits of table. Correlations point the right way only on average, not on every draw. In this case, Werner got the much less informative branch, and paid the premium price for it. The destroyed-per-received ratio of this round is 0.589: well below 1, anti-leverage.

Question 1 is a reminder that the forecast rows are *expectations*: they hold on average over many repetitions of the experiment, and only if the world really does draw its truth ψ from Werner's prior. On any single run the realized numbers are free to deviate.

EXAMPLE

Question 2. The state after the collapse, with fresh forecasts:

$a \backslash q$	11	10	01	00
0	0	0.855	0.831	0.658
1	1	0.145	0.169	0.342
$\langle s \rangle$	—	0.596	0.655	0.927
$\langle \Delta H(M) \rangle$	—	0.670	0.677	0.983



Greedy picks $q = 00$. Nature answers $a = 0$ — the answer Werner favored (probability 0.658), so the surprisal is only $s_2 = 0.604$ bits. Four of the eight die:

j	$\phi_j(00)$	name	p_j	p'_j	
8	0	$b_1 \wedge b_0$ (AND)	0.513	0.780	
9	1	$b_1 \equiv b_0$ (XNOR)	0.222		
10	0	b_0 (projection)	0.096	0.147	
11	1	$b_1 \rightarrow b_0$ (IF THEN)	0.024		
12	0	b_1 (projection)	0.024	0.037	
13	1	$b_0 \rightarrow b_1$ (IF THEN)	0.072		
14	0	$b_1 \vee b_0$ (OR)	0.024	0.037	
15	1	TRUE (constant 1)	0.024		

The ledger:

$$H(M): 2.178 \rightarrow 1.065, \quad H(p): 2.093 \rightarrow 1.035, \quad C: 0.085 \rightarrow 0.030.$$

This time 1.113 bits are destroyed for 0.604 received. Where did the extra 0.509 come from? Partly from the favored answer arriving — the belief dropped 1.058 against only 0.604 paid — and partly from the store: C fell by 0.055. Step ratio 1.842. Cumulatively, however, 1.860 destroyed per 1.873 received gives a leverage of 0.993 after 2 steps: still a hair below 1. Even a strong second round has not quite repaid the expensive shock of question 1.

Why track C , when its step-by-step changes visibly do not match the amounts deduced? Because on a single realized run the deduced part has *two* sources, and C is the only quantity that separates them. Writing Δ for the realized drops in one step,

$$\underbrace{\Delta H(M) - s}_{\text{deduced}} = \underbrace{[\Delta H(p) - s]}_{\text{windfall}} + \underbrace{[-\Delta C]}_{\text{tower}}, \quad (32)$$


an identity, since $H(M) = H(p) + C$ at every stage. The windfall is luck: the realized surprisal undershooting (or overshooting) the realized drop in belief entropy: zero in expectation. The

tower term is structure: a genuine withdrawal from the correlation store, and the channel with a global guarantee. Over a complete run its withdrawals sum to exactly the initial C , while the windfalls average away.

EXAMPLE

Question 3. Two questions remain:

$a \backslash q$	11	10	01	00
0	0	0.927	0.817	1
1	1	0.073	0.183	0
$\langle s \rangle$	—	0.378	0.687	—
$\langle \Delta H(M) \rangle$	—	0.408	0.717	—



Greedy picks $q = 01$; nature answers $a = 0$ (probability 0.817, surprisal $s_3 = 0.292$). Two die:

j	$\phi_j(01)$	name	p_j	p'_j	
8	0	$b_1 \wedge b_0$ (AND)	0.780	0.955	
10	1	b_0 (projection)	0.147		
12	0	b_1 (projection)	0.037	0.045	
14	1	$b_1 \vee b_0$ (OR)	0.037		

The ledger:


$$H(M): 1.065 \rightarrow 0.264, \quad H(p): 1.035 \rightarrow 0.264, \quad C: 0.030 \rightarrow 0.$$

Again a bargain: 0.801 destroyed for a mere 0.292 received, and the cumulative ratio finally clears unity, at 1.229. And note that C has dropped to zero.

C measures correlations that are not visible from M . If there's only one unknown column, there's nothing to infer. Put another way, usually, the binary maps are an overcomplete basis for M , because their number $N \gg \dim(M) = 2^{m+n}$. That's not the case when the number of free parameters has been reduced by learning to 1 each.

Question 4. One question, one heavily loaded coin:

$a \backslash q$	11	10	01	00
0	0	0.955	1	1
1	1	0.045	0	0
$\langle s \rangle$	—	0.264	—	—
$\langle \Delta H(M) \rangle$	—	0.264	—	—



The two forecasts now coincide: with $C = 0$ there is no transfer left to hope for, and the last question can only be observed, not leveraged. Nature answers $a = 0$ (probability 0.955, surprisal $s_4 = 0.066$):

j	$\phi_j(10)$	name	p_j	p'_j
8	0	$b_1 \wedge b_0$ (AND)	0.955	1.000
12	1	b_1 (projection)	0.045	—

The ledger closes:

$$H(M): 0.264 \rightarrow 0, \quad H(p): 0.264 \rightarrow 0, \quad C: 0 \rightarrow 0.$$

The truth is AND. The step destroyed 0.264 for 0.066 received (ratio 3.990) — not deduction this time, but luck: the destroyed amount was fixed in advance (the column's full entropy), while the realized surprisal undershot it because the likelier answer arrived. In the split of equation (32): pure windfall, tower 0.

Cumulatively: 2.925 bits of table destroyed, all of it, for 2.231 bits of surprisal received: a run ratio of 1.311. Most of the realized 1.311 was luck, not intelligence. Figure 4 stacks the books after each question: what remains, plus what was received. Once it dips below the starting entropy, our prior has done some work.




Figure 4: The realized run through the guiding example: remaining table entropy plus cumulative surprisal received, after each question, against the initial total (dashed). The gap between each stack and the line is the cumulative deduced part. At $\ell = 1$ the stack tops out 0.521 bits *above* the line — the negative deduction of question 1, plainly visible. At $\ell = 2$ the stack still grazes the line from above (0.013 bits); only from $\ell = 3$ on do the stacks fall short of it, as genuine deduction overtakes the early loss.

2.4 Leverage

If the prior is good, meaning it actual models reality (the prevalence of maps), we expect graphs such as the one above to drop as questions come in. The ratio of table entropy destroyed per bit of surprisal received deserves a name.

The **leverage**. Run the protocol for k rounds: questions $q_1, \dots, q_k$ are asked, the answers $a_i = \psi(q_i)$ are each processed by the update rule (27), producing the tables $M^{(i)}$, and round i costs the surprisal $s_i = -\log_2 M_{a_i, q_i}^{(i-1)}$. The cumulative leverage after k rounds is the table entropy destroyed per bit of surprisal received:

$$L_k := \frac{H(M^{(0)}) - H(M^{(k)})}{\sum_{i=1}^k s_i}. \quad (33)$$

A leverage of 1 is the uniform-prior baseline: every bit of table entropy destroyed was paid for at face value, one bit of surprisal each. Anything above 1 is deduction — table entropy removed without ever being observed.

2.5 Correlators

This section can be skipped on a first reading: the later chapters do not depend on it.

What kind of operator is the update rule (27)? In map space it is as simple as an operator gets: multiplication by a diagonal 0/1 mask, then renormalization. But there is a change of basis in which the same operation becomes a *convolution* with a two-term kernel, and in which the correlation store C stops being a single number and becomes visible coordinate by coordinate. The main advantage is that the effect of each update is well tabulated beforehand into a small number of targeted parameters.

This section builds that basis for $m = 1$; the final paragraph deals with general m .

First, re-encode each answer as a spin,

$$\sigma_q(\phi_j) := (-1)^{\phi_j(q)} \in \{+1, -1\}, \quad (34)$$

so answer 0 $\mapsto +1$ and 1 $\mapsto -1$. A pure relabeling, but it buys the one algebraic fact everything below turns on: $\sigma_q^2 = 1$.

The **Walsh–Hadamard transform** is built from a single matrix. For one bit,

$$W = \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}, \quad W_{c,b} = (-1)^{cb}, \quad c, b \in \{0, 1\}. \quad (35)$$

EXAMPLE

Acting on a one-bit distribution,

$$W \begin{pmatrix} p_0 \\ p_1 \end{pmatrix} = \begin{pmatrix} p_0 + p_1 \\ p_0 - p_1 \end{pmatrix} = \begin{pmatrix} 1 \\ \langle \sigma \rangle \end{pmatrix} :$$

row $c = 0$ computes the normalization, row $c = 1$ the spin bias. Distribution in,

correlators out.

The normalization riding along is a redundancy. The original distribution also uses one variable more than its degrees of freedom.

EXAMPLE

Tensoring two copies (the $n = 1$ case) does the same for two bits at once:

$$W \otimes W = \begin{pmatrix} W & W \\ W & -W \end{pmatrix} = \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & -1 \\ 1 & 1 & -1 & -1 \\ 1 & -1 & -1 & 1 \end{pmatrix} \quad \begin{array}{l} \leftarrow c = 00: \text{ sums to 1} \\ \leftarrow c = 01: \langle \sigma_0 \rangle \\ \leftarrow c = 10: \langle \sigma_1 \rangle \\ \leftarrow c = 11: \langle \sigma_1 \sigma_0 \rangle \end{array} \quad (36)$$

The Kronecker product multiplies signs, so for any number l of tensor factors the entry at row c , column j is

$$(W^{\otimes l})_{c,j} = \prod_{b=0}^{l-1} (-1)^{c_b j_b} = (-1)^{\sum_b c_b j_b}, \quad (37)$$

with c_b, j_b the b -th bits of the row and column indices; every further tensor factor appends one more bit.

For $m = 1$ a map *is* its truth table of Q bits — the numeral j — so Werner's belief is a vector of length $N = 2^Q$, and one transform computes every correlator at once:

$$\chi = W^{\otimes Q} p, \quad \chi_{\mathcal{S}} = \left\langle \prod_{q \in \mathcal{S}} \sigma_q \right\rangle = \sum_{j=0}^{N-1} (-1)^{\sum_{q \in \mathcal{S}} \phi_j(q)} p_j. \quad (38)$$

Each row computes the **correlator** of one subset $\mathcal{S} \subseteq \mathcal{Q}$ of questions. In the subscript we may encode the set as its *mask*: written in the same descending question order as everywhere else, bit q of the subscript records whether $q \in \mathcal{S}$. The row index of the matrix, read as a numeral, is the mask.

EXAMPLE

For $n = 2$, the question order is $(11, 10, 01, 00)$ so χ_{0110} is the correlator of $\mathcal{S} = \{10, 01\}$.

the sixteen rows of $W^{\otimes 4}$ probe the sixteen subsets of $\mathcal{Q} = \{11, 10, 01, 00\}$. A few entries of the dictionary:

mask	set $\mathcal{S}$	correlator
0000	$\emptyset$	$\chi_{0000} = 1$ (normalization)
0001	$\{00\}$	$\chi_{0001} = \langle \sigma_{00} \rangle$
1000	$\{11\}$	$\chi_{1000} = \langle \sigma_{11} \rangle$
0110	$\{10, 01\}$	$\chi_{0110} = \langle \sigma_{10} \sigma_{01} \rangle$
1111	$\{11, 10, 01, 00\}$	$\chi_{1111} = \langle \sigma_{11} \sigma_{10} \sigma_{01} \sigma_{00} \rangle$, the full parity

Sort the N correlators by $|\mathcal{S}|$, the number of questions probed, into *shells*. Shell 0 is the normalization, $\chi_{\emptyset} = 1$, riding along as a correlator. Shell 1 is M in different clothes:

$$\chi_{\{q\}} = M_{0,q} - M_{1,q}, \quad M_{a,q} = \frac{1}{2} (1 + (-1)^a \chi_{\{q\}}), \quad (39)$$


the bias of column q . For $m = 1$, we elegantly see the first shell correlators are the distance from fair coins.

Shell 2 holds the agreement biases of pairs of questions — knowledge M cannot see: two priors can share every column of M and differ here. Higher shells store subtler and subtler parity knowledge. Nothing is lost in the sorting: $W^2 = 2I$, so $(W^{\otimes Q})^2 = NI$ — the transform is its own inverse up to a factor N , and knowing every correlator is knowing the belief. The ledger also gives C a signature: a product prior ($C = 0$) factorizes every correlator, $\chi_S = \prod_{q \in S} \chi_{\{q\}}$, so its upper shells repeat what shell 1 already said; correlated knowledge is precisely *non-factorizing* correlators.

EXAMPLE

The prior of Table 4 transforms to the ledger

shell	correlators
0	$\chi_{0000} = 1$
1	$\chi_{1000} = +0.170, \chi_{0100} = +0.558, \chi_{0010} = +0.780, \chi_{0001} = +0.636$
2	$\chi_{1100} = -0.032, \chi_{1010} = +0.230, \chi_{1001} = +0.374,$ $\chi_{0110} = +0.498, \chi_{0101} = +0.434, \chi_{0011} = +0.576$
3	$\chi_{1110} = +0.028, \chi_{1101} = +0.092, \chi_{1011} = +0.434, \chi_{0111} = +0.374$
4	$\chi_{1111} = +0.152$



Shell 1 reproduces the columns of equation (26): $(1 + 0.170)/2 = 0.585 = M_{0,11}$, and likewise 0.779, 0.890, 0.818. The negative entry χ_{1100} says questions 11 and 10 disagree slightly more often than they agree; every other pair leans toward agreement. This is a reflection of the strong coherent 0-preference encoded in the prior. Answering 1 would give negative 1-shells.

Now look at the *rows* of $W^{\otimes Q}$ as functions of j . They are **block waves**. The observation mask is built of the same material:

$$\delta(\phi_j(q), a) = \frac{1}{2} \left(1 + (-1)^{a+\phi_j(q)} \right), \quad (40)$$

a 0/1 block wave — one constant plus one Hadamard row. And multiplication and convolution swap under this transform exactly as under Fourier: multiplying two functions of j pointwise convolves their ledgers. And the observation mask's ledger has only two nonzero entries — at $\emptyset$ and at $\{q\}$ — so this convolution is no great sum: it combines each correlator with exactly one other.

The update rule, carried out entirely in the ledger, is therefore that two-term combination, followed by renormalization by the new shell 0. In index notation, learning $a = \psi(q)$ sends every set S to

$$\chi'_S = \frac{\chi_S + (-1)^a \chi_{S \triangle \{q\}}}{1 + (-1)^a \chi_{\{q\}}}, \quad (41)$$

where $S \triangle \{q\}$ is S with the membership of q toggled: q is added if absent, removed if present (on masks, bit q is flipped). The numerator is the monomial algebra of $\sigma_q^2 = 1$: multiplying

$\prod_{r \in S} \sigma_r$ by σ_q toggles q 's membership of the product, so every correlator is averaged with its **partner**: the correlator differing by q , signed by the answer. The denominator is the $S = \emptyset$ instance of the numerator, and by equation (39) it equals $2M_{a,q}$: the same number that priced the answer in equation (27) renormalizes the ledger.

EXAMPLE




Figure 5: The update rule in the ledger, drawn for the guiding example's first question ($q = 11$). Every set is paired with the partner whose membership of 11 is toggled (arrows); each pair is averaged with the answer's sign and renormalized, per equation (41). The pair $\emptyset \leftrightarrow \{11\}$ at the orange-bordered asked set fixes the normalization and saturates the asked bias; each blue double arrow averages a pair into both of its members, and the net flow of structure is one shell down, toward M . The thin gray arrows from the asked q to every other pill remind that this sets the renormalization. The fill encodes each prior value to scale: darker is larger, orange is negative.

Read formula (41) by cases:

- $S = \emptyset$: partner $\{q\}$, giving $(1 + (-1)^a \chi_{\{q\}}) / (1 + (-1)^a \chi_{\{q\}}) = 1$. Normalization is preserved.
- $S = \{q\}$, the asked question: partner $\emptyset$, giving $(\chi_{\{q\}} + (-1)^a) / (1 + (-1)^a \chi_{\{q\}}) = (-1)^a$. The bias saturates to certainty — the one-hot collapse of column q , derived rather than decreed.
- $S = \{q'\}$, an unasked question: the partner is the pair $\{q', q\}$, and one line of algebra puts the shift in terms of the *connected* correlator:

$$\chi'_{\{q'\}} - \chi_{\{q'\}} = \frac{(-1)^a (\chi_{\{q',q\}} - \chi_{\{q'\}} \chi_{\{q\}})}{2M_{a,q}} = \frac{(-1)^a \text{Cov}(\sigma_{q'}, \sigma_q)}{2M_{a,q}}. \quad (42)$$


Unasked columns move in proportion to their covariance with the asked question: knowledge stored one shell up descends into the predictions. A product prior has zero covariance everywhere, so its unasked columns are frozen — the “ $C = 0$ means no transfer” fact, now a one-line consequence.

- Higher shells: the same cascade. Each update marries shell- k sets containing q to shell- $(k-1)$ sets without it; whatever structure was stored jointly with q moves one shell down, toward M .

EXAMPLE

Question 1 of the walkthrough asked $q = 11$ and received $a = 1$: $\text{sign}(-1)^a = -1$, toggled set $\{11\}$, mask 1000. Equation (41) maps the ledger of the previous box to

shell	updated correlators
0	$\chi'_{0000} = 1$
1	$\chi'_{1000} = -1$, $\chi'_{0100} = +0.711$, $\chi'_{0010} = +0.663$, $\chi'_{0001} = +0.316$
2	$\chi'_{1100} = -0.711$, $\chi'_{1010} = -0.663$, $\chi'_{1001} = -0.316$, $\chi'_{0110} = +0.566$, $\chi'_{0101} = +0.412$, $\chi'_{0011} = +0.171$
3	$\chi'_{1110} = -0.566$, $\chi'_{1101} = -0.412$, $\chi'_{1011} = -0.171$, $\chi'_{0111} = +0.267$
4	$\chi'_{1111} = -0.267$



The asked bias saturates: $\chi'_{1000} = (0.170 - 1)/(1 - 0.170) = -1$. The other singletons move by equation (42): for $q' = 10$ the covariance is $-0.032 - (0.558)(0.170) = -0.127$, so the shift is $(-1)(-0.127)/(2 \times 0.415) = +0.153$, landing on $+0.711$ — and indeed $(1 + 0.711)/2 = 0.855$, the column $M_{0,10}$ of the Question 2 box. Note also that every mask containing the asked bit is now just $(-1)^a$ times its partner ($\chi'_{1010} = -\chi'_{0010}$, etc.): with σ_{11} pinned, no correlator involving question 11 carries independent knowledge anymore.





Figure 6: The update rule of equation (41), drawn for the guiding example's first update ($q = 11$, $a = 1$, so $(-1)^a = -1$ and toggled bit 1000). One row per set $\mathcal{S}$. At $x = 2M_{a,q} = 1 + (-1)^a \chi_{1000} = 0.830$ stands the numerator as a waterfall: the bar of $\chi_{\mathcal{S}}$ (blue), and beside it the partner's signed contribution $(-1)^a \chi_{\mathcal{S} \Delta \{q\}}$ (orange), hanging from the level where the first bar ends. The two subscripts differ by the flipped bit. The ray from the origin through the waterfall's endpoint rescales it, by similar triangles, to $x = 1$, where its height is the updated correlator $\chi'_{\mathcal{S}}$.

As Werner moves through the questions, more correlators settle to ± 1 . The 1-shell models the belief of the bias of each question. The 2-shell answers questions such as 'If the answer to $\psi(10) = 0$, then $\psi(01) = 1$ '. And the higher orders are even more conditional, measuring belief about such statements, in the case of certain answers, etc. As we learn more, the conditions become satisfied or not, and the information becomes less nested.

EXAMPLE

The remaining questions of the run (00, 01, 10, all answered 0, so with sign +1) each saturate their own bit the same way, and since XOR toggles commute, the order does not matter — the commutativity of updating, seen again. After all four, the ledger is pure signs, $\chi_{\mathcal{S}} = (-1)^{|\mathcal{S} \cap \{11\}|}$ — -1 on every set containing 11, $+1$ on the rest: the transform of the point mass on $\psi = \phi_8$:

shell	final correlators
0	$\chi_{0000} = 1$
1	$\chi_{1000} = -1$, $\chi_{0100} = +1$, $\chi_{0010} = +1$, $\chi_{0001} = +1$
2	$\chi_{1100} = -1$, $\chi_{1010} = -1$, $\chi_{1001} = -1$, $\chi_{0110} = +1$, $\chi_{0101} = +1$, $\chi_{0011} = +1$
3	$\chi_{1110} = -1$, $\chi_{1101} = -1$, $\chi_{1011} = -1$, $\chi_{0111} = +1$
4	$\chi_{1111} = -1$



2.5.1 General m

It turns out that taking $m > 1$ is more general but not much more instructive. With multiple output bits it is as if we have m a priori independent maps per question; they just happen always to be asked and answered simultaneously, in unison. Therefore we do not have to think about what one of them teaches us about another. It would mean we have more than 2 options every question, and therefore the factor by which every update thins the hypotheses is not 2 but A . The guiding example helps because the maps are represented as binary numbers, which are more familiar.

For completeness, the main formulas. The truth table is now mQ bits: one *site* (q, i) per bit i of each answer, each carrying a spin $\sigma_{q,i}(\phi_j) := (-1)^{\phi_j(q)_i}$, where $\phi_j(q)_i$ is the i -th

bit of the answer. Correlators are indexed by sets $\mathcal{S}$ of sites — masks of mQ bits under the same numeral convention, m of them per question:

$$\chi_{\mathcal{S}} = \left\langle \prod_{(q,i) \in \mathcal{S}} \sigma_{q,i} \right\rangle = \sum_{j=0}^{N-1} (-1)^{\sum_{(q,i) \in \mathcal{S}} \phi_j(q)_i} p_j. \quad (43)$$

EXAMPLE

For $n = 3$, $m = 2$ the sites form a 2×8 grid: one column per question, one row per bit of the answer. A set $\mathcal{S}$ is a pattern of hot cells in this grid, and $\chi_{\mathcal{S}}$ multiplies the corresponding spins:

$q =$	11	110	101	100	011	010	001	000
bit 1								
bit 0								
	$\chi_{\emptyset} = 1$: the normalization							
bit 1								
bit 0								
	$\langle \sigma_{101,1} \rangle$: the bias of one answer bit							
bit 1								
bit 0								
	$\langle \sigma_{110,1} \sigma_{110,0} \rangle$: the parity of one answer — with the two biases, the block of $M_{\cdot,110}$							
bit 1								
bit 0								
	$\langle \sigma_{100,0} \sigma_{011,1} \rangle$: agreement across two questions							
bit 1								
bit 0								
	$\langle \sigma_{111,1} \sigma_{101,0} \sigma_{010,1} \sigma_{010,0} \sigma_{000,0} \rangle$: a five-site parity, deep in the shells							

Any of the $2^{16} = 65536$ hot/cold patterns is a correlator, and sorting by the number of hot cells gives the shells. The mask of $\mathcal{S}$ is the grid read as a 16-bit numeral.

Write $\mathcal{B}_q$ for the block of q 's m sites, and $a \cdot \mathcal{T} := \sum_{i \in \mathcal{T}} a_i$ for the parity of the answer a on a subset $\mathcal{T} \subseteq \mathcal{B}_q$. The answer matrix is read off the correlators supported on one block — A of them per column:

$$M_{a,q} = \frac{1}{A} \sum_{\mathcal{T} \subseteq \mathcal{B}_q} (-1)^{a \cdot \mathcal{T}} \chi_{\mathcal{T}}. \quad (44)$$

The observation mask at (q, a) is a function of the m bits of one block, so its ledger has A nonzero entries, and hearing $a = \psi(q)$ combines every correlator with its A partners — one per toggle $\mathcal{T}$ within the asked block:

$$\chi'_{\mathcal{S}} = \frac{\sum_{\mathcal{T} \subseteq \mathcal{B}_q} (-1)^{a \cdot \mathcal{T}} \chi_{\mathcal{S} \Delta \mathcal{T}}}{A M_{a,q}}, \quad (45)$$

the denominator once more the $\mathcal{S} = \emptyset$ instance of the numerator. For $m = 1$ the three formulas collapse to equations (38), (39) and (41).

EXAMPLE

To draw the update, take a sparse prior at $n = 3$, $m = 2$: four maps, $\phi(q) = q \bmod 4$ (weight 0.45), constant 00 (0.35), “11 iff $q \geq 4$ ” (0.12), constant 11 (0.08). Werner asks $q = 110$ and hears $a = 11$: toggle signs $(+, -, -, +)$ on $(\emptyset, 110_0, 110_1, \text{both})$, where q_i denotes site (q, i) , and denominator $AM_{11,110} = 0.8$. The pairs of the $m = 1$ figure

become *quartets*:




Figure 7: For $m = 2$, toggling any subset of the asked block's two sites ties four correlators into a clique that averages jointly, per equation (45). The on-block quartet (orange hull) contains the normalization and the entire block of $M_{\cdot,110}$; its signed sum is the denominator that renormalizes every other quartet. Fills to scale for the sparse prior, as before.




Figure 8: The $m = 2$ update, three sets from the sparse prior. The waterfall now has four segments: χ_S (blue) and its three signed quartet partners, in toggle order 110_0 , 110_1 , both (light to dark orange). The ray from the origin rescales the endpoint from $x = AM_{a,q} = 0.8$ to $x = 1$. The asked site saturates to -1 ; the foreign site 011_0 flips sign, $-0.060 \rightarrow +0.200$; and $\{011_0, 110_1\}$ lands on -0.200 , the slaved value $-\chi'_{\{011_0\}}$.

3 Expectations

In expectation over the truth $\psi \sim p$, as well as over the question order, we take the ratio of the expected totals,

$$\langle L_\ell \rangle := \frac{H(M^{(0)}) - \langle H(M^{(\ell)}) \rangle}{\langle \sum_{i=1}^{\ell} s_i \rangle}. \quad (46)$$

We did this implicitly in Section 1.2.1. There, the problem was actually symmetric in $q_0 \leftrightarrow q_1$. That is not the case in general.

In order to progress, we will be interested in such averages. Both are natural. The question order is uniform, without replacement: Without control of the questions, Werner has no reason to expect one to arrive before another. And the truth is drawn from the prior itself: Werner scores his future against his own beliefs, and he has faith in them.

The bookkeeping simplifies at once. Updating commutes, so after ℓ rounds the state depends only on the asked *set* $\mathcal{S}$, a uniform random subset of size ℓ . Let us say the answers

received form a vector $y = \{y_i\}$. The surprisals telescope to the surprisal of the whole answer block,

$$\sum_{i=1}^{\ell} s_i = -\log_2 \Pr(\psi(\mathcal{S}) = y), \quad (47)$$

where we remember the pairing from $\mathcal{S} \ni q_i \mapsto y_i$ implicitly. The average over ψ is a block entropy. Everything is therefore generated by one sequence, the mean block entropy at size k :

$$G_k := \binom{Q}{k}^{-1} \sum_{|\mathcal{S}|=k} H(\psi(\mathcal{S})), \quad G_0 = 0, \quad G_Q = H(p). \quad (48)$$

G_1 is the mean column entropy, so $H(M)_0 = QG_1$. Two exact laws follow. The first is equation (47) averaged:

$$\left\langle \sum_{i=1}^{\ell} s_i \right\rangle = G_{\ell}. \quad (49)$$

For the second, fix $\mathcal{S}$ and average the posterior table over the answers. Each unasked q' column's entropy becomes a conditional entropy, $H(\psi(q') \mid \psi(\mathcal{S})) = H(\psi(\mathcal{S} \cup \{q'\})) - H(\psi(\mathcal{S}))$, and the asked columns contribute zero. Now average over $\mathcal{S}$: every set of size $\ell+1$ arises as $\mathcal{S} \cup \{q'\}$ once per element, and counting with $(\ell+1)\binom{Q}{\ell+1} = (Q-\ell)\binom{Q}{\ell}$,

$$\langle H(M^{(\ell)}) \rangle = (Q-\ell)(G_{\ell+1} - G_{\ell}). \quad (50)$$


The whole expected learning curve is the discrete derivative of one monotone sequence.

EXAMPLE

Take the guiding space, $\mathcal{Q} = \{11, 10, 01, 00\}$ with $Q = 4$ and $m = 1$, and abbreviate $\psi(11, 10)$ for $\psi(\{11, 10\})$. Definition (48) averages $H(\psi(\mathcal{S}))$ over the sets of each size, so the five block entropies collect 1, 4, 6, 4, 1 terms apiece:

$$\begin{aligned} G_0 &= 0, \\ G_1 &= \frac{1}{4} [H(\psi(11)) + H(\psi(10)) + H(\psi(01)) + H(\psi(00))], \\ G_2 &= \frac{1}{6} [H(\psi(11, 10)) + H(\psi(11, 01)) + H(\psi(11, 00)) \\ &\quad + H(\psi(10, 01)) + H(\psi(10, 00)) + H(\psi(01, 00))], \\ G_3 &= \frac{1}{4} [H(\psi(10, 01, 00)) + H(\psi(11, 01, 00)) \\ &\quad + H(\psi(11, 10, 00)) + H(\psi(11, 10, 01))], \\ G_4 &= H(\psi(11, 10, 01, 00)) = H(p). \end{aligned}$$

Those sixteen sets are the nodes of a lattice, one rank per value of ℓ , and each G_{ℓ} is the average of the entropies on rank ℓ :



Putting together (49) and (50), equation (46) becomes

$$\langle L_\ell \rangle = \frac{H(M^{(0)}) - (Q - \ell)(G_{\ell+1} - G_\ell)}{G_\ell}. \quad (51)$$

In terms of the prior, the block entropies are explicit. The maps that answer the pattern y on the asked set $\mathcal{S}$ form a residue class, with mass

$$G_\ell = -\binom{Q}{\ell}^{-1} \sum_{|\mathcal{S}|=\ell} \sum_y P_{\mathcal{S},y} \log_2 P_{\mathcal{S},y}, \quad P_{\mathcal{S},y} := \sum_{j: \phi_j(q_i)=y_i \forall q_i \in \mathcal{S}} p_j, \quad (52)$$

the inner sum running over the A^ℓ answer patterns y on the block $\mathcal{S}$. Nothing here is specific to $m = 1$.

Some observations.

- Expectation kills the windfall. Each step has $\langle \Delta H(p) - s \rangle = 0$, so the expected deduced part is a pure tower withdrawal: $H(M^{(0)}) - \langle H(M^{(\ell)}) \rangle - G_\ell = C_0 - \langle C_\ell \rangle$, with $\langle C_\ell \rangle = \langle H(M^{(\ell)}) \rangle - (H(p) - G_\ell)$. It is nonnegative: anti-leverage is only ever bad luck.
- Deduction saturates one round early: substituting $G_Q = H(p)$ at $\ell = Q - 1$ already gives $\langle C_{Q-1} \rangle = 0$. The store is empty before the last question, which can only be observed, never leveraged.
- $\langle L_\ell \rangle \geq 1$ always, but it need not be monotone. The endpoint is forced: $\langle L_Q \rangle = H(M^{(0)})/H(p) = 1 + C_0/H(p)$, the run ratio quoted in Section 2. This is also why equation (46) is a ratio of expectations: the expectation of the ratio is dominated by lucky runs with vanishing denominators.
- Sanity: $\ell = 0$ returns $H(M^{(0)}) = QG_1$, and at $n = 1$, half-support prior, the laws reproduce the one-parameter world of Section 1.2.1.

There is also a monotonicity theorem hiding in G . Entropy is submodular, so the increments $G_\ell - G_{\ell-1}$, which are the expected surprisals of the ℓ -th answer, decrease with ℓ : answers get cheaper as evidence accumulates. This of course only holds on uniformly random question order.

EXAMPLE

For the guiding prior the block-entropy sequence is $G_k = (0, 0.731, 1.436, 2.100, 2.707)$, and the laws unfold it into the expected run:

	received	remaining	store	leverage
ℓ	$\langle \sum_i s_i \rangle$	$\langle H(M^{(\ell)}) \rangle$	$\langle C_\ell \rangle$	$\langle L_\ell \rangle$
1	0.731	2.113	0.137	1.110
2	1.436	1.328	0.056	1.113
3	2.100	0.608	0	1.104
4	2.707	0	0	1.081

The expected leverage humps at $\ell = 2$ and lands on $2.925/2.707 = 1.081$. The realized run of Section 2 scattered around this gentle expectation: greedy questioning and a lucky truth delivered 1.311. Figure 9 stacks the expected books.




Figure 9: The expected run through the guiding example. Conservation is now exact at every ℓ : remaining plus received plus deduced reaches $H(M)_0$, so the cake always tops out on the dashed line, and no overshoot is possible. The top layer is the expected withdrawal from the correlation store, $C_0 - \langle C_\ell \rangle$; it saturates at $C_0 = 0.218$ already at $\ell = 3$.

3.1 Choosing the questions

The uniform order was a symmetry choice, not a strategy. Werner can do better along two independent axes: how far he looks ahead, and whether he commits to an order beforehand or lets each question depend on the answers so far. *Greedy* asks, at every step, the single question with the largest expected drop of $H(M)$: one step of lookahead. ℓ -*optimal* maximizes the expected knowledge after exactly ℓ questions, looking the whole horizon ahead. All four combinations are straightforward with the machinery at hand.

	order fixed beforehand	contingent on answers
greedy	Grow the asked set one question at a time, each maximizing the expected drop of $H(M)$ averaged over the still-unseen answers. One pass of the bookkeeping of equation (52); front-loads the cheap destruction, but cannot react to luck.	Recompute the forecasts of Section 2 after every answer and ask the current argmax: the walkthrough’s play. One update per step; beats its fixed twin whenever an answer reshuffles the ranking.
ℓ -optimal	Updating commutes, so a fixed plan is just a set: search the $\binom{Q}{\ell}$ subsets for the smallest expected remainder $\sum_{q' \notin \mathcal{S}} H(\psi(q') \mid \psi(\mathcal{S}))$. The best non-adaptive knowledge; sees correlations that the one-step view misses.	Backward induction over posteriors, $V_t = \min_q \mathbb{E}_a V_{t-1}$ with $V_0 = H(M)$. The Bayes-optimal policy: dearest to compute, and the ceiling the other three approximate.

Expected knowledge improves left to right and top to bottom. The dynamic program in the last cell is exact but grows quickly: its states are the reachable posteriors, one per pair of asked set and answer pattern, $(1 + A)^Q$ in all. That is 81 for the guiding example and astronomical beyond; its policy is also horizon-dependent, since the best first question for $\ell = 1$ need not open the best plan for $\ell = 2$. For the guiding prior, contingent greedy already attains the contingent optimum at every horizon: expected remaining 1.821, 0.903, 0.334 bits, against the uniform 2.113, 1.328, 0.608. That is a kindness of this prior, not a law: a prior that stores its knowledge in deep shells can defeat greedy in either column, since a question worth little now may unlock a cascade later.

3.2 The Gibbs complexity prior

Every prior so far was engineered: a mechanism first, a distribution second. This section builds the prior the other way around, from a single principle – simple explanations are likelier – with circuit complexity as the measure of simplicity, made quantitative as a Gibbs ensemble over the hypothesis space. It is an illustration: the following chapter does not depend on it.

3.2.1 NAND complexity, and two cheaper descriptors

Fix the universal gate: the two-input NAND. The complexity X_j of a map ϕ_j is the fewest gates in any feed-forward circuit that computes it. For the bookkeeping we count in the equivalent *and-inverter* form: two-input AND gates of unit cost, with inversion free on every wire – an edge attribute, not a gate. De Morgan’s laws push inverters along wires and into gates, so an and-inverter circuit converts to a NAND-only circuit of the same gate count and back; X_j is the standard AIG-size metric of logic synthesis. Reference values: AND, OR, NAND and NOR all cost 1; XOR costs 3; constants, projections and negations cost 0. For $m > 1$ outputs, X_j is the size of one *joint* circuit computing all output bits together, gates shared – the honest minimum-hardware cost, smaller in general than the sum of per-output circuits.

The tabulation is feasible because X_j is invariant under relabeling: permuting or negating inputs, negating or (for $m = 2$) swapping outputs are all free operations, so the 65,536 maps

of the $(4, 1)$ and $(3, 2)$ spaces collapse to 222 and 308 equivalence classes. The $(4, 1)$ class values are published optimal AIG sizes¹; the $(3, 2)$ values come from that synthesizer: circuit sizes $r = 0, 1, 2, \dots$ are encoded as satisfiability instances, and the first satisfiable r is provably minimal, every smaller size having been refuted.

Alongside the interior cost X_j , which takes a SAT solver to see, two descriptors are visible in one pass over the truth table:

- the *footprint* $F_j = n \cdot [\text{Image}] + [\text{Support}]$, where the support counts the input bits the map actually depends on and the image size counts the distinct answers it attains (from 1 for constants up to A). The combination is lossless – base- n digits, quotient I , remainder S – and reads as the map’s *interface*: does it react to all its inputs, does it use all its values?
- the *bias* B_j : the number of 1s minus the number of 0s among the mQ truth-table bits. This is the symmetry breaker: X_j and F_j are even under negating all outputs, B_j is odd – an energy term coupling to it acts like a magnetic field, switching on the odd correlator shells that a pure complexity prior forbids.

EXAMPLE

For the guiding $(2, 1)$ space, with each map’s prior mass under the Occam and Pragmatist field settings of Section 3.2.2 drawn as opposed bars about a shared axis – Occam in blue growing left, the Pragmatist in orange growing right, full bar width one unit of probability:

j	map	X_j	Support	Image	F_j	B_j	prior mass	
							Occam	Pragmatist
0	FALSE	0	0	1	2	-4	0.62	0.28
1	$\neg(b_1 \vee b_0)$ (NOR)	1	2	2	6	-2	0.11	0.00
2	$\neg b_1 \wedge b_0$	1	2	2	6	-2	0.11	0.00
3	$\neg b_1$	0	1	2	5	0	0.11	0.00
4	$b_1 \wedge \neg b_0$	1	2	2	6	-2	0.11	0.00
5	$\neg b_0$	0	1	2	5	0	0.11	0.00
6	$b_1 \oplus b_0$ (XOR)	3	2	2	6	0	0.11	0.00
7	$\neg(b_1 \wedge b_0)$ (NAND)	1	2	2	6	+2	0.11	0.00
8	$b_1 \wedge b_0$ (AND)	1	2	2	6	-2	0.11	0.00
9	$b_1 = b_0$ (XNOR)	3	2	2	6	0	0.11	0.00
10	b_0 (projection)	0	1	2	5	0	0.11	0.00
11	$b_1 \rightarrow b_0$	1	2	2	6	+2	0.11	0.00
12	b_1 (projection)	0	1	2	5	0	0.11	0.00
13	$b_0 \rightarrow b_1$	1	2	2	6	+2	0.11	0.00
14	$b_1 \vee b_0$ (OR)	1	2	2	6	+2	0.11	0.00
15	TRUE	0	0	1	2	+4	0.28	0.62

Occam spreads its belief across the six gate-free maps (0.1106 apiece) and thins gradually with cost; the Pragmatist, punishing footprint and rewarding ones, piles 0.90 of its mass onto the two constants – 0.62 on TRUE against 0.28 on FALSE, the tilt breaking the tie – and all but abandons the interior of the table. Grouped by the full tuple (X, F, B) , the sixteen maps fall into six energy classes:

¹The github.com/krinkin/bounds reproducibility dataset, cross-checked against our own SAT-based synthesizer on random classes with zero mismatches.

(X, F, B)	members	(X, F, B)	members
$(0, 2, -4)$	1	$(1, 6, -2)$	4
$(0, 2, +4)$	1	$(1, 6, +2)$	4
$(0, 5, 0)$	4	$(3, 6, 0)$	2

Note the guiding space skips $X = 2$ entirely: with two inputs, nothing a two-gate circuit computes is out of reach of one gate or zero.

At the larger sizes the tuple breakdown runs to hundreds of rows (102 distinct tuples at $(4, 1)$, 196 at $(3, 2)$), so we tabulate the complexity tiers alone, with each tier’s *aggregate* prior mass under the Occam setting alongside its head count:

EXAMPLE

$(4, 1)$			$(3, 2)$		
X	maps	Occam mass	X	maps	Occam mass
0	10	0.041 	0	64	0.058 
1	48	0.072 	1	432	0.144 
2	256	0.141 	2	2,064	0.253 
3	940	0.190 	3	4,860	0.219 
4	2,336	0.174 	4	10,832	0.180 
5	6,464	0.177 	5	16,208	0.099 
6	10,616	0.107 	6	17,060	0.038 
7	18,984	0.070 	7	10,672	0.009 
8	17,680	0.024 	8	3,312	0.001 
9	7,882	0.004 	9	32	0.000 
10	320	0.000 			
total	65,536	1	total	65,536	1

The bars are magnified twofold relative to the butterfly table.

The head counts are thin at both ends: ten gate-free maps at $(4, 1)$ (two constants, eight literals) against only 320 at the ceiling $X = 10$. The aggregate mass shows the thermodynamic compromise this forces on Occam: the count grows faster than e^{-X} decays until $X = 3$ at $(4, 1)$ ($X = 2$ at $(3, 2)$), so even under a pure simplicity prior the *typical* map has middling complexity — tier degeneracy beats tier energy until well past the simplest maps.

3.2.2 The ensemble

Take the energy linear in the three descriptors and the prior as its Boltzmann weight,

$$E_j = \beta X_j + \alpha F_j + \mu B_j, \quad p_j = \frac{e^{-E_j}}{\sum_k e^{-E_k}}, \quad (53)$$

with no overall temperature: the couplings are the fields. Positive β suppresses complex maps, positive α suppresses large footprints, and μ couples to the signed bias, $\mu < 0$ tilting toward 1-heavy tables and thereby breaking the output-negation symmetry. Three settings are explored below:

	β	α	μ
Occam	1	0	0
Occamer	2	0	0
Pragmatist	1.5	1	-0.1

Occam suppresses circuit complexity and nothing else. *Occamer* doubles the same field: an even stronger simplicity bias stores more knowledge in the prior, and stores it deeper in the correlation shells, where only the later questions of a run can release it. The *Pragmatist* judges by the obvious heuristics instead: a full footprint field reads the visible interface rather than the interior circuit, a moderate complexity field seconds it, and the tilt $\mu = -0.1$ leans toward 1-heavy tables, breaking the output-negation symmetry that the other two settings keep.

3.2.3 Trajectories

Figures 10–12 show the expected cumulative leverage $\langle L_\ell \rangle$ under the three settings, for the uniform-order average and the strategy grid of Section 3.1. The ℓ -optimal schedule targets the halfway horizon $\ell^* = Q/2$ and is plotted only up to it: past its own horizon the optimizer’s continuation is not defined, and inside the optimal *set* the questions are asked in random order. Its contingent twin is omitted everywhere: wherever we can afford to compute it, it reproduces the fixed optimum’s curve to machine precision in all three settings – a coincidence of these fields rather than a law (a stronger tilt, $\mu = -0.3$, separates the two), and one for which we have no symmetry argument yet (at $(4, 1)$ its 3^{16} -state program is in any case out of reach). Two anchors hold in every panel: all schedules agree at $\ell = 1$, because input relabeling acts transitively on the questions and the Gibbs energies are relabeling-invariant, so single questions are exchangeable; and every full-length curve ends at the same $L_Q = H(M)_0/H(p)$, since any policy that asks everything receives exactly $H(p)$.




Figure 10: Expected leverage under the Gibbs priors at $(2, 1)$, $\ell^* = 2$. In both complexity-only settings every schedule coincides exactly; the Pragmatist’s fields finally separate them.




Figure 11: Expected leverage under the Gibbs priors at $(3, 2)$, $\ell^* = 4$. Greedy's fixed set is the optimal set under Occam and, up to an input negation, under Occamer; under the Pragmatist it misses – knowing less at the horizon yet out-levering the optimum.




Figure 12: Expected leverage under the Gibbs priors at $(4, 1)$, $\ell^* = 8$. Occamer's deep store bends even the strategyless average curve into an interior hump; the Pragmatist is a two-sided spike on the constants, all leverage and little game.

Three observations, one per system size.

At $(2, 1)$, under *both* complexity-only settings, the strategy grid collapses exactly: all question subsets of equal size share a single joint entropy, so every schedule is the uniform order. Occamer shows what doubling the field buys within that straitjacket – the whole curve lifts (endpoint 1.2224 against Occam's 1.0829, store $C_0 = 0.728$ against 0.306 bits) without creating a single strategic choice to make. The Pragmatist's fields split the pair orbits at last, and the very first thing strategy does is post a *lower* number: greedy's pair $\{00, 11\}$ knows more at $\ell = 2$ than the random order and shows less leverage for it (2.538 against 2.582).

At $(3, 2)$, the optimal four questions are the even-weight inputs $\{000, 011, 101, 110\}$ – the maximally spread design, pairwise at Hamming distance two – under *all three* settings. Under Occam greedy builds exactly that set; under Occamer it lands on the odd-weight twin, the same design up to an input negation, and its curve coincides with the optimum's. Under the Pragmatist greedy finally misses: its set $\{000, 011, 101, 111\}$ leaves twice the ignorance of the optimal four (0.405 against 0.198 bits remaining) while receiving 0.099 bits less – and it posts the *higher* leverage, 4.445 against 4.355, with the strategyless average higher still at 4.476. The ℓ -optimal schedule maximizes knowledge, not leverage, and a ratio is perfectly happy with cheaper ignorance.

At $(4, 1)$ the settings differ far more than the schedules within any one of them, and

Occamer is the showpiece: doubling β triples the store ($C_0 = 8.10$ bits against Occam's 2.48) and bends even the average curve into an interior hump, rising to 2.20 at $\ell = 6$ before decaying – the deep-shell signature promised in Section 3.2.2, knowledge that only becomes spendable once enough answers are in hand to cash the higher correlations. The optimal eight questions are a parity class of the input hypercube in every setting (even-weight under Occam and the Pragmatist, the odd twin under Occamer); greedy misses the class under both complexity-only settings (1.249 against the optimum's 1.262 at the Occam horizon) but finds it exactly under the Pragmatist – whose prior, 0.90 of its mass on the two constants, is a two-sided spike in disguise: $H(p) = 0.53$ bits, leverage running from 12.29 down to 8.28, a great deal of inference about a game that is mostly over before it starts.

4 The thermodynamic limit

The earlier chapters describe learning on a finite question set. This chapter asks what remains when the number of input bits tends to infinity. The hope is to find traces of universality: that character of learning where the microscopic details average out.

Since $Q = |\mathcal{Q}| = 2^n$, the limit $n \rightarrow \infty$ is one of an exponentially growing table, and we measure time by the fraction of that table already queried,

$$t = \frac{\ell}{Q} \in [0, 1]. \quad (54)$$

The limit at fixed $t > 0$ is *macroscopic*: it records learning that occupies a nonzero fraction of the run. Learning completed in $O(1)$, in $O(\log Q)$, or more generally in $o(Q)$ questions is compressed into a boundary layer at $t = 0$. That distinction will determine the class of solutions we find.

Throughout, $u := A^{-1} = 2^{-m}$ abbreviates the reciprocal answer alphabet, and $N = A^Q$ counts the maps $\mathcal{Q} \rightarrow \mathcal{A}$, enumerated $\phi_0, \dots, \phi_{N-1}$ with prior weights $p_j = \Pr(\psi = \phi_j)$. One value of the truth ψ is drawn at the start of a run and supplies every answer in it; posteriors are conditional distributions of that same random map, and expected trajectories average over its prior law.

4.1 A toy example: the spike prior and its renormalization flow

The spike prior places a distinguished weight on one map and spreads the remaining mass uniformly over every other map. Take the special map to be the all-zero map, $\phi_0(q) = 0^m$ for every $q \in \mathcal{Q}$, so that its answer to any k questions is the all-zero block 0^{mk} . This is fully general by symmetry: the answer alphabet may be relabelled separately at each question. Then

$$p_0 = p, \quad p_j = \omega := \frac{1-p}{N-1} \quad (j \neq 0). \quad (55)$$

The plain $p = p_0$ is the mass of the all-zero map alone. We will be interested in the behavior at large (but not necessarily unit) p . The interpretation is that Werner is quite sure of his hypothesis, without ruling anything out.

The example is useful because it is closed under Bayesian conditioning. A confirming answer produces a sharper spike on a smaller map space; a refuting answer removes the special map and leaves equal weights on all survivors. The posterior therefore has a two-state flow, in which the uniform family is the stable, absorbing class: once a trajectory enters it, later conditioning keeps it there. An interior spike is a source of one-way probability current into that class, while conditional on avoiding refutation the surviving spike sharpens.

4.1.1 The two-state flow and exact leverage

What happens when Werner handles the first question? He either gets a zero, or not.

The maps that produce 0^m are a proportion $u = A^{-1}$. One has probability p , and $Nu - 1$ have probability ω . If we started with probability $P_0 = 1$ in this prior, Werner avoids refuting ϕ_0 with probability

$$P_1 = p + (Nu - 1)\omega \quad (56)$$

In this case, what exactly is the shape of the posterior? The Bayesian consistency condition tells us that unconditional probability cannot be created without measurement. At this point in our ansatz, we must have the same credence in ϕ_0 . Let r_1 be the height of the peak p_0 , after one question. Then we must have $r_1 \cdot P_1 = p$. The rest of the probability is spread again over the $Nu - 1$ other surviving maps. Conversely, with probability $1 - P_1$, we've refuted the zero map, and are in a uniform posterior corresponding to a nonzero answer to the question.

After repeating this process k times, asking the set of questions $\mathcal{S}$, we can infer the probability that ϕ_0 is still in the race. Let P_k be the probability that we remain in the confirming scenario: Werner has only seen zeros. A history remains on the spike branch exactly when $\psi(\mathcal{S}) = 0^{mk}$, and exactly Nu^k maps produce that confirming history, one of which is ϕ_0 . Its prior probability is therefore

$$P_k := \Pr(\psi(\mathcal{S}) = 0^{mk}) = p + (Nu^k - 1)\omega. \quad (57)$$

On that event the posterior is another spike on those Nu^k maps,

$$r_k = \Pr(\psi = \phi_0 \mid \psi(\mathcal{S}) = 0^{mk}) = \frac{p}{P_k}, \quad \omega_k = \frac{\omega}{P_k}, \quad (58)$$

so r_k is the time-indexed posterior height of the spike evaluated on the realized confirming branch, with $r_0 = p$ the original weight. It's pegged inversely to P_k for all time, by the same Bayesian consistency:


$$P_k r_k = p. \quad (59)$$

If any answer is nonzero then ϕ_0 is eliminated, every compatible survivor had the same mass ω_k , and the posterior is uniform; subsequent conditioning preserves uniformity. With an equivalence across system sizes, this induces a two state flow. Renormalization sharpens the spike. The two macrostates and their directions are

$$\text{spike} \longrightarrow \begin{cases} \text{sharper spike,} & \text{with probability } P_{k+1}/P_k, \\ \text{uniform,} & \text{with probability } 1 - P_{k+1}/P_k, \end{cases} \quad \text{uniform} \longrightarrow \text{uniform.} \quad (60)$$

EXAMPLE

The flow at $(n, m) = (3, 2)$, so $A = 4$ and $N = 4^8 = 65,536$, started from $r_0 = p_0 = 0.3$. Each confirmed answer divides the surviving map space by four and renormalizes the spike upward; each refutation drops the belief into the uniform class, which nothing ever leaves.



For $p \in (0, 1)$, we can combine equations (55) and (57) to get the limiting odds on the surviving branch. Also the surviving posterior, from equation (58):

$$P_k = p + (1 - p)u^k + O(N^{-1}), \quad r_k = \frac{p}{p + (1 - p)u^k} + O(N^{-1}). \quad (61)$$

Each confirmation supplies approximately $-\log_2 u = m$ bits of evidence for the spike. After Q (all) questions, the only way the spike can survive is by eliminating all alternative maps. This happens with $P_Q = p$ and hence $r_Q = 1$; the histories that confirm through the complete table are precisely those with $\psi = \phi_0$. At the uniform point $p = 1/N$ both posterior branches are uniform on their surviving map spaces.

The same P_k has a second reading: it is the probability of the single k -answer block 0^{mk} , so there is no new distribution to compute. For $1 \leq k \leq Q$ every other answer block is realized by Nu^k background maps and therefore carries the common probability

$$U_k := \frac{1 - P_k}{A^k - 1} = Nu^k \omega, \quad \text{so} \quad \Pr(\psi(\mathcal{S}) = y) = \begin{cases} P_k, & y = 0^{mk}, \\ U_k, & y \neq 0^{mk}, \end{cases} \quad (62)$$

for any question set $\mathcal{S}$ of size k . Note that $A = u^{-1}$. The entropy of the above block distribution is

$$G_k = -P_k \log_2 P_k - (A^k - 1) U_k \log_2 U_k. \quad (63)$$

Every choice of k questions obeys the same law, so (63) is also the mean k -question block entropy of Chapter 3. In particular $G_0 = 0$, $G_Q = H(\psi)$ and $H(M)_0 = Q G_1$, with the exact

prior entropy

$$G_Q = h_2(p) + (1-p) \log_2(N-1), \quad (64)$$

extensive in Q at fixed $0 < p < 1$ and vanishing at $p = 1$.

4.1.2 The fixed- p thermodynamic limit

The thermodynamic limit keeps the prior mass $0 < p < 1$ fixed and sends $n \rightarrow \infty$ at a fixed asked fraction $t > 0$. The guiding principle is the exact finite law (51): the expected leverage is built from three ingredients,

1. the prior table entropy $H(M^{(0)})$
2. the received entropy G_ℓ
3. the increment $G_{\ell+1} - G_\ell$.

If we can derive each for the example, we are done. The block entropy (63) supplies every G_k through the single counting formula (57), while $H(M^{(0)}) = Q G_1$ holds for any prior: G_1 is by definition the mean entropy of one column. Set $\ell = \lfloor tQ \rfloor$ and take the terms of (51) in turn.

All three derivations run through one exact regrouping of (63). The $A^k - 1$ refuting blocks (62) share the mass $1 - P_k$ evenly, so substituting $U_k = (1 - P_k)/(A^k - 1)$ into (63) and expanding the logarithm regroupes the block entropy into the binary entropy of the confirm-or-refute split plus a uniform choice among the refuting blocks:

$$G_k = h_2(P_k) + (1 - P_k) \log_2(A^k - 1). \quad (65)$$

The prior table entropy: G_1 is the one-column entropy, (65) at $k = 1$. By (57) the spike answer 0^m carries $P_1 = p + (Nu - 1)\omega \rightarrow p + (1 - p)u$: the special map's whole mass, plus the fraction u of the background maps that happen to agree with it at any one question. Substituting that limit,

$$\begin{aligned} \frac{H(M^{(0)})}{Q} &= G_1 = h_2(P_1) + (1 - P_1) \log_2(A - 1) \\ &\rightarrow h_2(p + (1 - p)u) + (1 - p)(1 - u) \log_2(A - 1) =: h_{\text{spike}}(p, m). \end{aligned} \quad (66)$$

The received entropy: Evaluate (65) at $k = \ell = \lfloor tQ \rfloor$. By (61) $P_\ell \rightarrow p$, so the split entropy stays bounded while the block-count logarithm grows linearly, $\log_2(A^\ell - 1) = \ell m + o(1)$:

$$\frac{G_{\lfloor tQ \rfloor}}{Q} = \underbrace{\frac{h_2(P_\ell)}{Q}}_{\rightarrow 0} + \underbrace{(1 - P_\ell)}_{\rightarrow 1 - p} \underbrace{\frac{\log_2(A^\ell - 1)}{Q}}_{\rightarrow tm} \rightarrow t(1 - p)m. \quad (67)$$

This is the two-state flow watched at macroscopic time: for any $t > 0$ the confirm-or-refute lottery has already resolved inside the boundary layer at $t = 0$, where $r_k \rightarrow 1$ within $o(Q)$ questions. A surviving spike answers for free; only the refuted class, probability $1 - p$, keeps paying the full m bits per answer.

The increment: Subtract consecutive copies of (65): the bounded $h_2(P_k)$ terms cancel in the limit, and the block-count logarithm grows by $\log_2(A^{\ell+1} - 1) - \log_2(A^\ell - 1) \rightarrow \log_2 A = m$, paid only by the refuted mass:

$$G_{\ell+1} - G_\ell \rightarrow (1 - p)m, \quad \text{so} \quad \frac{(Q - \ell)(G_{\ell+1} - G_\ell)}{Q} \rightarrow (1 - t)(1 - p)m. \quad (68)$$

Filling the three limits into (51) term by term,

$$\langle L_\ell \rangle = \frac{\overbrace{Q h_{\text{spike}}(p, m)}^{H(M^{(0)})} - \overbrace{(Q - \ell)(G_{\ell+1} - G_\ell)}^{G_\ell}}{\underbrace{Q t(1 - p)m}_{\rightarrow Q t(1 - p)m}},$$

and cancelling the common factor Q gives

$$L_p(t) = \frac{h_{\text{spike}}(p, m) - (1 - t)(1 - p)m}{t(1 - p)m} = 1 + \frac{h_{\text{spike}}(p, m) - (1 - p)m}{t(1 - p)m}, \quad 0 < t \leq 1. \quad (69)$$

The $1/t$ term records information released in a vanishing initial fraction of the run, divided by the $O(tQ)$ information received afterwards.




Figure 13: The fixed- p thermodynamic curve (69), on a logarithmic vertical scale. Every weight produces the same $1/t$ shape; the prior sets only the coefficient, which grows from 0.001 at $p = 0.001$ to above 5 at $p = 0.999$.

We have explored a number of other characteristics of the spike prior in Appendix A: the branch-by-branch effect of the first question, the extraction ratio, and various further limits, among them the sharp prior $p \rightarrow 1^-$ and its noncommutation with $n \rightarrow \infty$.

4.2 Thermodynamic-limit calculus

The spike is solvable because all subsets of the same size have the same entropy. The general theory replaces that symmetry by an average over subsets.

4.2.1 The finite sequence from which the limit is taken

For each fixed question set $\mathcal{S} \subseteq \mathcal{Q}$ the vector $\psi(\mathcal{S})$ collects the answers to those questions, and its entropy is taken under the prior law of ψ with $\mathcal{S}$ held fixed. The mean block entropy is that of equation (48), now carrying the system size explicitly:

$$G_{n,k} := \binom{Q}{k}^{-1} \sum_{\substack{\mathcal{S} \subseteq \mathcal{Q} \\ |\mathcal{S}|=k}} H(\psi(\mathcal{S})), \quad 0 \leq k \leq Q, \quad (70)$$

so that $G_{n,0} = 0$ and $G_{n,Q} = H(\psi) = H(p^{(n)})$, where $p^{(n)}$ is the prior law of the random map at size n , with increments $\gamma_{n,k} := G_{n,k+1} - G_{n,k}$.

By the entropy chain rule, exactly as in the derivation of equation (50), the increment is the mean conditional entropy of one fresh answer,

$$\gamma_{n,k} = \mathbb{E}_{\mathcal{S},q}[H(\psi(q) \mid \psi(\mathcal{S}))], \quad |\mathcal{S}| = k, \quad q \in \mathcal{Q} \setminus \mathcal{S}. \quad (71)$$

There are two separate averages here: the conditional entropy averages over answer histories under the prior on the random map, while $\mathbb{E}_{\mathcal{S},q}$ averages that entropy over which questions were selected. Conditioning reduces entropy, so coupling a random k -set to a random $(k+1)$ -set yields

$$m \geq \gamma_{n,0} \geq \gamma_{n,1} \geq \cdots \geq \gamma_{n,Q-1} \geq 0. \quad (72)$$

This monotonicity is deterministic, the truth and the subset having already been averaged over. It should not be called self-averaging; concentration of individual learning trajectories would be a separate result. The two laws (49) and (50) then read

$$\left\langle \sum_{i=1}^{\ell} s_i \right\rangle = G_{n,\ell}, \quad \langle H(M^{(\ell)}) \rangle = (Q - \ell) \gamma_{n,\ell}, \quad (73)$$

for $0 \leq \ell < Q$, with zero remaining entropy at $\ell = Q$, and equation (51) becomes

$$\langle L_{n,\ell} \rangle = \frac{Q G_{n,1} - (Q - \ell) \gamma_{n,\ell}}{G_{n,\ell}} \quad (1 \leq \ell < Q), \quad \langle L_{n,Q} \rangle = \frac{Q G_{n,1}}{G_{n,Q}}. \quad (74)$$

4.2.2 A precise macroscopic-limit statement

It is useful to keep the initial marginal entropy separate from the macroscopic profile on the right, since that allows an $o(Q)$ -question boundary layer at $t = 0$. Let $h_0 := \lim_{n \rightarrow \infty} G_{n,1}$, and suppose the staircases have an almost-everywhere limit γ on $(0, 1)$, with $\gamma_{n, \lfloor tQ \rfloor} \rightarrow \gamma(t)$ at any particular t where the remaining entropy is evaluated. At such a point $\gamma(t)$ is the expected entropy of one more answer after a random t -fraction has been observed. Four hypotheses suffice:

1. **Fixed answer width.** The integer m is fixed, which supplies the uniform bound $0 \leq \gamma_{n,k} \leq m$.
2. **Initial-entropy convergence.** The limit h_0 exists.
3. **Profile convergence.** The staircase $t \mapsto \gamma_{n, \lfloor tQ \rfloor}$ converges almost everywhere on $(0, 1)$, with direct convergence at any t where the remaining entropy is evaluated. At continuity points of a selected monotone limit this causes no ambiguity.
4. **Positive entropy density.** The limiting area is positive, $\int_0^1 \gamma(x) dx > 0$.

The fourth condition is the nondegenerate form of extensivity, and it implies $G_{n,Q}/Q \rightarrow \int_0^1 \gamma > 0$. Merely writing $G_{n,Q} = \Theta(Q)$ is not sufficient: the entropy density may oscillate with n , and even a convergent scalar density does not determine the profile. Conversely, profile convergence with positive area implies the extensivity required. No Gibbs representation and no full-support lower bound on individual map probabilities is needed.

Extend the finite increments to a staircase on $[0, 1]$ by $\gamma_n(t) := \gamma_{n, \min(\lfloor tQ \rfloor, Q-1)}$, the value at the single endpoint $t = 1$ being irrelevant to the integral. Telescoping gives an exact Riemann-sum identity at grid points,

$$\frac{G_{n,\ell}}{Q} = \frac{1}{Q} \sum_{k=0}^{\ell-1} \gamma_{n,k} = \int_0^{\ell/Q} \gamma_n(x) dx, \quad (75)$$

and since $0 \leq \gamma_n \leq m$, dominated convergence gives $G_{n,\ell_n}/Q \rightarrow g(t) := \int_0^t \gamma$ for $\ell_n = \lfloor tQ \rfloor$, while at a point of profile convergence $(Q - \ell_n)\gamma_{n,\ell_n}/Q \rightarrow (1-t)\gamma(t)$. Substituting both into (74) proves

$$L(t) = \frac{h_0 - (1-t)\gamma(t)}{g(t)}, \quad g(t) = \int_0^t \gamma(x) dx, \quad 0 < t < 1. \quad (76)$$

Since γ is nonnegative and nonincreasing, a positive integral makes $g(t) > 0$ for every $t > 0$.




Figure 14: From the finite increments to the profile, drawn for parity blocks of size four, whose limit is $\gamma(t) = 1 - t^3$ (dashed). Each bar has width $1/Q$ and height $\gamma_{n,k}$, so the shaded bars sum – the telescoping identity (75) – to exactly $G_{n,\ell}/Q$, here at $\ell/Q = 5/8$. Refining Q shrinks the bars onto the staircase’s limit, and dominated convergence turns the shaded sum into the area $g(t) = \int_0^t \gamma$: received information per question becomes the area under the profile.

Monotonicity and boundedness guarantee convergent subsequences by a Helly selection argument, but they do not guarantee that the full sequence of priors selects a unique profile: a family can alternate between two constructions on even and odd n . At the prior level one therefore needs a coherent construction, such as a fixed exchangeable directing measure or fixed local block types with convergent frequencies, or must take profile convergence itself as a hypothesis.

4.3 A constructive complexity prior

The Gibbs prior of Section 3.2 priced maps by circuit complexity: a principled Occam, but one whose complexities had to be tabulated by force, map by map, and which therefore stops at small n . This section builds an Occam prior from a complexity notion that is *constructive* – computable from the truth table by a fast transform – and tractable at every n : the degree of the map written as a polynomial. Two ideas carry the section. First, a non-hyperbolic macroscopic curve requires a non-exchangeable prior (de Finetti flattens any consistent exchangeable family), and the polynomial structure of the question cube supplies the asymmetry without any partition of the questions into blocks. Second, weighting complexity by a single temperature fails for a counting reason made precise below: only a thin shell of maps would contribute. The repair is to *sample* the complexity level rather than price it, which keeps every level – and every single map – at positive weight all the way into the thermodynamic limit. Throughout this section the answer width is $m = 1$, so $u = \frac{1}{2}$ and a fresh fair answer carries one bit.

4.3.1 Polynomials over $\mathbb{F}_2$, and the butterfly

Write $\mathbb{F}_2$ for the two-element field $\{0, 1\}$, in which addition is XOR ($1 \oplus 1 = 0$) and multiplication is AND; all linear algebra in this section is over $\mathbb{F}_2$. Since a bit obeys $b^2 = b$, no variable is ever squared, and a monomial in the input bits $b_{n-1} \cdots b_0$ of a question is labelled by the *set* S of bits it contains. Every map is a XOR of monomials, and the correspondence is one-to-one: there are 2^n monomials, hence $2^Q = N$ coefficient vectors, exactly as many as maps. This is the *algebraic normal form*,

$$\phi(q) = \bigoplus_S a_S \prod_{i \in S} b_i, \quad a_S = \bigoplus_{y: \text{supp}(y) \subseteq S} \phi(y), \quad (77)$$

with single-bit coefficients $a_S \in \{0, 1\}$: monomial S is present or absent. The second formula is Möbius inversion: each coefficient is the XOR of the truth table over the subcube below S . The *degree* $\deg \phi$ is the largest $|S|$ with $a_S = 1$; it is the complexity notion of this section.

EXAMPLE

OR answers 1 everywhere except at $q = 00$. Its four coefficients, by (77): $a_\emptyset = \phi(00) = 0$; $a_{\{b_0\}} = \phi(00) \oplus \phi(01) = 1$; $a_{\{b_1\}} = \phi(00) \oplus \phi(10) = 1$; $a_{\{b_1 b_0\}} = \phi(00) \oplus \phi(01) \oplus \phi(10) \oplus \phi(11) = 1$. Hence $\text{OR} = b_1 b_0 \oplus b_1 \oplus b_0$, of degree 2: OR is not affine, and its truth table knows it.

Mechanically the inversion is a butterfly of XORs. Per bit it is an upper-triangular kernel, and on n bits its tensor power is the subset indicator; in the numeral convention of Table 4, questions descending 11, 10, 01, 00 (per bit: 1, 0),

$$Z = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}, \quad (Z^{\otimes n})_{S,y} = \mathbf{1}[y \subseteq S], \quad Z^{\otimes 2} = \begin{pmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 \end{pmatrix}, \quad (78)$$

where $\mathbf{1}[\cdot]$ evaluates a statement to a bit. Over $\mathbb{F}_2$ the kernel is its own inverse, $Z^2 = I$. In place: for each bit i , for every x with $x_i = 1$, do $T[x] \oplus = T[x \text{ with bit } i \text{ cleared}]$; $n 2^{n-1}$ XORs convert a truth table to its coefficients. For readers of Section 2.5: Z is the triangular half of the Walsh–Hadamard kernel (35), since in the same 1, 0 ordering

$$W = \begin{pmatrix} -1 & 1 \\ 1 & 1 \end{pmatrix} = Z \begin{pmatrix} -2 & 0 \\ 0 & 1 \end{pmatrix} Z^\top, \quad (79)$$

the matrix form of $\prod_q (1 - 2y_q) = \sum_T (-2)^{|T|} \prod_{q \in T} y_q$: the ledger transform is two Möbius passes with a $(-2)^{|S|}$ rescaling between them, and reducing mod 2 leaves exactly the XOR butterfly.

EXAMPLE

The butterfly on six maps of Table 4, truth tables $T = (\phi(11), \phi(10), \phi(01), \phi(00))$ – so that T read as a binary numeral *is* the index j – and coefficients ordered $(a_{b_1 b_0}, a_{b_1}, a_{b_0}, a_\emptyset)$:

map (j)	T	a	polynomial	deg
TRUE (15)	(1, 1, 1, 1)	(0, 0, 0, 1)	1	0
$\neg b_1$ (3)	(0, 0, 1, 1)	(0, 1, 0, 1)	$b_1 \oplus 1$	1
XOR (6)	(0, 1, 1, 0)	(0, 1, 1, 0)	$b_1 \oplus b_0$	1
AND (8)	(1, 0, 0, 0)	(1, 0, 0, 0)	$b_1 b_0$	2
NAND (7)	(0, 1, 1, 1)	(1, 0, 0, 1)	$b_1 b_0 \oplus 1$	2
$b_1 \rightarrow b_0$ (11)	(1, 0, 1, 1)	(1, 1, 0, 1)	$b_1 b_0 \oplus b_1 \oplus 1$	2

AND's table equals its own coefficient vector (a single top monomial); NAND flips only a_0 , the algebraic face of negation; XOR and $\neg b_1$ are affine; TRUE is constant.

4.3.2 Complexity classes and the sampling rule

Grade the hypothesis space by degree. The class of all maps of degree at most d is a linear space of dimension $K_d = \sum_{i \leq d} \binom{n}{i}$ – one coefficient bit per allowed monomial – and the classes are nested, from the two constants ($d = 0$) to the full map space ($d = n$). In coding theory this class is the Reed–Muller code $\text{RM}(d, n)$. Its *rate* is the dimension per question,

$$R_{n,d} = \frac{K_d}{Q} \in (0, 1], \quad (80)$$

and it will turn out to be the fraction of a full run at which the class is learned out. Consecutive rates slice $(0, 1]$ into *rate gaps* $(R_{n,d-1}, R_{n,d}]$ of lengths $\binom{n}{d}/2^n \leq \sqrt{2/\pi n}$: one cell per class, a partition that refines as n grows.

Why sample a level rather than weight the maps? The natural Occam move, $p_j \propto e^{-\beta \deg \phi_j}$, fails by a scaling argument. There are 2^{K_k} maps of degree at most k , so raising the cutoff by one buys $\binom{n}{k+1} \ln 2$ nats of entropy against a linear price β ; no single temperature balances jumps of that size at more than one k . The ensemble concentrates on the shell $\binom{n}{d^*} \approx \beta / \ln 2$ – at most two adjacent degrees – and every map outside it retains only vanishing relative weight. A single temperature hears a thin shell of the hypothesis space and nothing else. The repair prices the *level*, not the map:

1. draw a single random number $R \in [0, 1]$ (the *statistic*) from a fixed law μ with CDF $F(x) = \Pr(R \leq x)$ and, when it exists, density $f = F'$;
2. let D be the class whose rate gap contains R ;
3. draw the truth uniformly from class D : each allowed coefficient by a fair coin.

In every expectation below the random object is this one draw R ; the $R_{n,d}$ are the deterministic grid it lands among. The class label D is latent – Werner never observes it – and is one value among $n + 1$, worth at most $\log_2(n + 1)$ bits. Integrating the lottery, and using nestedness (a map of degree e belongs to every class $d \geq e$),

$$w_d = F(R_{n,d}) - F(R_{n,d-1}), \quad p_j = \sum_{d \geq \deg \phi_j} w_d 2^{-K_d}, \quad (81)$$

with $\sum_d w_d = 1$ telescoping. If F increases across every gap (any positive density), then every $w_d > 0$, and two properties follow at once: p_j is *strictly decreasing in the degree*, each step down costing a factor of order $2^{-\binom{n}{d}}$ – a genuine Occam – and $p_j \geq w_n 2^{-Q} > 0$ for every single map: full support, with no noise channel, at every n and in the limit. Every

level, and every map, gets a chance. Semantically, F budgets the *shells* while nestedness orders the *maps*: w_d is, up to a $2^{-\binom{n}{d}}$ sliver leaking in from higher classes, the total weight of the degree- d shell, split evenly within it. The per-map Occam ordering therefore holds for every F , and F decides only how much aggregate weight each complexity level carries – a populous shell can outweigh a simpler one even though each of its members is individually lighter.

4.3.3 Three terms for the leverage average

The exact finite law (51), in the chapter’s notation (74), needs three ingredients: the prior table entropy $Q G_{n,1}$, the received entropy $G_{n,\ell}$, and the increment $\gamma_{n,\ell}$. Each reduces to a statement about F .

The table entropy. A class is a linear code; make the structure explicit once. The *generator matrix* $\mathbf{G}$ of class d (sans serif, distinct from the block entropies $G_{n,\ell}$) is the $K_d \times Q$ matrix over $\mathbb{F}_2$ whose rows are the allowed monomials evaluated at every question, $\mathbf{G}_{S,q} = \prod_{i \in S} b_i(q)$, so that a uniform class member is $\psi = u^\top \mathbf{G}$ with u uniform on $\{0,1\}^{K_d}$. Every column of $\mathbf{G}$ is nonzero (its constant-monomial entry is 1), so each single answer is the image of a uniform u under a surjective linear map onto $\mathbb{F}_2$: a fair bit, in every class, and a mixture of fair bits is a fair bit. At every finite n , with no limit taken and no reference to μ ,

$$G_{n,1} = 1, \quad H(M^{(0)}) = Q. \quad (82)$$

The prior predicts nothing about any single answer; its entire content is correlational.

The received entropy. Computing $H(\psi(\mathcal{S}), D)$ by the chain rule in both orders,

$$H(\psi(\mathcal{S})) = \sum_d w_d H_d(\psi(\mathcal{S})) + I(D; \psi(\mathcal{S})), \quad 0 \leq I \leq \log_2(n+1) : \quad (83)$$

mixture entropy is mean class entropy plus what the answers reveal about the label, and the label is cheap – $o(1)$ bits per question. Within class d the answers on $\mathcal{S}$ are $u^\top \mathbf{G}_{\cdot, \mathcal{S}}$, uniform on the image of the column submatrix, so $H_d(\psi(\mathcal{S})) = \text{rank}(\mathbf{G}_{\cdot, \mathcal{S}})$ exactly: the number of $\mathbb{F}_2$ -linearly independent columns, the number of bits needed to describe those answers. Three facts pin the limit of the class average $G_{n,\ell}^{(d)} = \mathbb{E} \text{rank}$. The mean increment $\gamma_{n,\ell}^{(d)} = \Pr[\text{a fresh column is independent of } \ell \text{ random ones}]$ is nonincreasing in ℓ , by rank submodularity (what is dependent on few columns is dependent on more). The increments telescope to the dimension, $\sum_{\ell < Q} \gamma_{n,\ell}^{(d)} = K_d$: the finite- n area theorem, received equals learnable. And the one imported result – Reed–Muller codes achieve capacity on the binary erasure channel (Kuddekar, Kumar, Mondelli, Pfister, Şaşoğlu and Urbanke, 2017) – says that revealing any fraction beyond the rate determines a fresh answer with probability tending to one:

$$R_{n,d_n} \rightarrow R \in (0,1) \implies \lim_{n \rightarrow \infty} \gamma_{n, \lfloor tQ \rfloor}^{(d_n)} = 0 \quad \text{for every fixed } t > R. \quad (84)$$

A bounded nonincreasing sequence with area R_d and no mass above R_d has nowhere to sit but at height 1 below, so the class increments converge to a pure step, written with the Heaviside function $\theta(x) = \mathbf{1}[x > 0]$:

$$\lim_{n \rightarrow \infty} \gamma_{n, \lfloor tQ \rfloor}^{(d)} = \theta(R_d - t), \quad \frac{G_{n, \lfloor tQ \rfloor}^{(d)}}{Q} \longrightarrow \int_0^t \theta(R_d - x) dx = \min(t, R_d) : \quad (85)$$

each class pays one full bit per question until its clock stops at its own rate, and nothing after. Finally, round the draw R up to the top of its rate gap: since $\min(t, \cdot)$ is 1-Lipschitz

and the rounding moves R by at most $\sqrt{2/\pi n}$, the class average converges to an average over the statistic,

$$\frac{G_{n, \lfloor tQ \rfloor}}{Q} \longrightarrow \mathbb{E}_\mu[\min(t, R)] = \int_0^t (1 - F(x)) dx =: g(t), \quad (86)$$

the last step by layer cake: pointwise $\min(t, R) = \int_0^t \mathbf{1}[x < R] dx$, and $\mathbb{E} \mathbf{1}[x < R] = 1 - F(x)$. Received information per question is an average of clipped clocks, the area under the survival function.

The increment. The sequence $\ell \mapsto G_{n, \ell}$ is concave (increments nonincreasing), and slopes of concave functions converge wherever the limit is differentiable, so at every continuity point of F

$$\gamma_{n, \lfloor tQ \rfloor} \longrightarrow g'(t) = 1 - F(t) = \Pr(R > t). \quad (87)$$

This is the heart of the construction: the mean conditional entropy of a fresh answer is the probability that the true class is not yet exhausted. The survival function of the statistic is not a postulate; it is the increment term of the finite law.

4.3.4 The macroscopic curve

Substitute (82), (86) and (87) into (51) and cancel the common factor Q ; the numerator tidies to $1 - (1 - t)(1 - F(t)) = t + (1 - t)F(t)$, and

$$L_F(t) = \frac{t + (1 - t)F(t)}{\int_0^t (1 - F(x)) dx}. \quad (88)$$

This is the master law (76) realized with $h_0 = 1$, $\gamma(t) = 1 - F(t)$ and $g(t) = \int_0^t (1 - F)$; the hypotheses hold for free, γ being bounded and nonincreasing by construction. The numerator has its own reading: destroyed density is the asked fraction t plus $F(t)$ per unasked column, since an unasked answer is deducible exactly when the class died before t . Asked plus deduced, over received. Expanding at the ends, $L_F(0^+) = 1 + f(0^+)$ and $L_F(1) = 1/\mathbb{E}_\mu[R]$: the curve starts at baseline plus the density of the statistic at zero, and completes at the reciprocal mean rate. Conversely, any nonincreasing profile with values in $[0, 1]$ is the survival function of exactly one law μ (read the layer cake backwards), so (88) is not one family among many: the sampling rule realizes every admissible curve of (76) at $m = 1$, and the statistic is the profile. A parallel construction from blocks of questions, mathematically complete in the same sense but less aligned with modeling a learner, is collected in Appendix B.

4.3.5 Two statistics: $F(x) = x$ and $F(x) = x^2$




Figure 15: The two worked statistics in (88). A uniform draw ($F(x) = x$, blue) gives constant leverage $L \equiv 2$; the larger of two uniform draws ($F(x) = x^2$, orange) delays the deductions and the curve rises from the baseline to $L(1) = \frac{3}{2}$.

A single uniform draw, $F(x) = x$. The numerator of (88) is $t + (1-t)t = t(2-t)$ and the denominator is $t - t^2/2 = t(2-t)/2$; the factor $t(2-t)$ cancels and $L \equiv 2$, exactly, at every macroscopic time. Both endpoint identities agree: $1 + f(0^+) = 2$ and $1/\mathbb{E}[R] = 2$. The interpretation is the flattest possible Occam: uniform F spreads its budget evenly over the *rate axis*, so each shell receives exactly the share of coefficient space its degree opens, $w_d = \binom{n}{d}/2^n$ – not an equal share per level – and every received bit drags along exactly one deduced bit, uniformly over the whole run. The same constant curve is produced by pairing questions into equality cliques (Appendix B), a completely different microscopic story: the flattest curve does not identify its microscopics.

The larger of two draws, $F(x) = x^2$. Both draws land below x with probability x^2 , so the weights (81) become gaps of squared rates, and the curve is

$$L(t) = \frac{t + (1-t)t^2}{t - t^3/3} = \frac{3(1+t-t^2)}{3-t^2}, \quad (89)$$

rising from $L(0^+) = 1 + f(0^+) = 1$ (the density vanishes at zero, so the curve leaves the baseline flat) through $L' \propto (1-t)(3-t) > 0$ to $L(1) = 1/\mathbb{E}[R] = \frac{3}{2}$. Keeping the larger draw tilts belief toward richer worlds: deductions are late-loaded, so early answers pay full price and the reward arrives toward the end of the run, a rising curve throughout.

EXAMPLE

The prior (81) at $(n, m) = (2, 1)$ with $F(x) = x^2$. The classes have $(K_0, K_1, K_2) = (1, 3, 4)$ and rates $(\frac{1}{4}, \frac{3}{4}, 1)$, so $w = (\frac{1}{16}, \frac{1}{2}, \frac{7}{16})$, and the sixteen maps of Table 4 are priced in three tiers by their degree: the two constants at $\frac{1}{16} \cdot \frac{1}{2} + \frac{1}{2} \cdot \frac{1}{8} + \frac{7}{16} \cdot \frac{1}{16} = \frac{31}{256}$, the six affine maps at $\frac{23}{256}$, and the eight curved maps at $\frac{7}{256}$; the check $2 \cdot 31 + 6 \cdot 23 + 8 \cdot 7 = 256$ closes the distribution.

Appendix A continues the spike-prior example of Section 4.1. Appendix B collects the block-prior construction and its Bernstein profile language, a block-based route to the same family of macroscopic curves.

A More characteristics of the spike prior

This appendix continues the toy example of Section 4.1: the anatomy of a single question, the fraction of the dependency stock it extracts, and the sharp-prior limit with its order-of-limits structure.

A.1 One question, branch by branch

For a single question abbreviate $P := P_1$ and $U := U_1$, so that $P + (A - 1)U = 1$: the zero answer has probability P , while each particular nonzero answer has probability U . Here *correct* means that the observed answer agrees with the spike map at this question; it does not assert that the entire spike map is the truth. The surprisals are

$$s_{\text{correct}} = -\log_2 P, \quad s_{\text{wrong}} = -\log_2 U, \quad \langle s \rangle = P s_{\text{correct}} + (A - 1)U s_{\text{wrong}} = G_1, \quad (90)$$

a particular wrong answer having probability U , not $1 - P$. Before asking, every one of the Q columns carries the same answer distribution $(P, U, \dots, U)$ and hence entropy G_1 , so $H_{\text{before}} = H(M)_0 = Q G_1$. Conditional on confirmation the next unasked answer has probabilities P_2/P for zero and U_2/P for each nonzero symbol, with entropy

$$H_c := -\frac{P_2}{P} \log_2 \frac{P_2}{P} - (A - 1) \frac{U_2}{P} \log_2 \frac{U_2}{P}, \quad H_{\text{after,correct}} = (Q - 1)H_c, \quad (91)$$

the asked column being fixed while all $Q - 1$ unasked columns share that conditional marginal. Refutation instead leaves equal posterior weights on all compatible maps, so every unasked answer is uniform on A symbols and $H_{\text{after,wrong}} = (Q - 1)m$. Weighting the two branches,

$$\langle H(M^{(1)}) \rangle = P H_{\text{after,correct}} + (A - 1)U H_{\text{after,wrong}} = (Q - 1)(G_2 - G_1), \quad (92)$$

the first equality ordinary branch averaging, the second the $\ell = 1$ case of (50).




Figure 16: One-question anatomy of the spike prior against the special map's mass p : the branch surprisals (90) (solid), the answer-matrix entropy before the question (dotted) and on each branch after it (91) (dashed). Both panels carry $Qm = 16$ bits of table entropy at the uniform point.

For $p < 1$ the branch leverages divide the vertical distance between the pre-question curve and each branch curve by that branch's own surprisal,

$$L_{\text{correct}} = \frac{QG_1 - (Q - 1)H_c}{-\log_2 P}, \quad L_{\text{wrong}} = \frac{QG_1 - (Q - 1)m}{-\log_2 U}, \quad (93)$$

whereas the expected leverage is the ratio of the two expected totals,

$$\langle L_1 \rangle = \frac{P[QG_1 - (Q-1)H_c] + (A-1)U[QG_1 - (Q-1)m]}{G_1} = \frac{QG_1 - (Q-1)(G_2 - G_1)}{G_1}. \quad (94)$$

The branch formulas use realized surprisals; the expected formula divides the expected entropy drop by the expected surprisal G_1 and is therefore *not* the probability-weighted average of the two branch ratios.




Figure 17: Branchwise (93) and expected (94) leverage of the first question. A refuting answer carries so much surprisal that its own ratio falls below the uniform baseline and eventually below zero, yet the expectation stays well above 1 because the confirming branch is both likely and cheap.

A.2 The extraction ratio

The total correlation $C := H(M)_0 - G_Q = QG_1 - G_Q$ is the information stored in dependencies among answers rather than visible in their separate marginals. Only the part of the expected entropy drop beyond the received entropy G_1 comes out of that store, so when $C > 0$ the share of it freed by a single question is

$$R_1 := \frac{\langle \Delta H(M) \rangle - G_1}{C} = \frac{(Q-1)(2G_1 - G_2)}{QG_1 - G_Q}. \quad (95)$$

The two ratios differ only in their denominator: R_1 measures against the whole dependency stock, leverage against the received surprisal, and

$$R_1 = \frac{G_1}{C} (\langle L_1 \rangle - 1), \quad \langle L_1 \rangle = 1 + \frac{C}{G_1} R_1. \quad (96)$$

Thus R_1 is a stock-extraction fraction while $\langle L_1 \rangle - 1$ is released dependency information per received bit. At $p = 1$ both the numerator of R_1 and C vanish; expanding (63) with $\varepsilon = 1 - p \rightarrow 0^+$ gives the exact finite- (n, m) ratio limit

$$\lim_{p \rightarrow 1^-} R_1(p) = \frac{(Q-1)N(1-u)^2}{QN(1-u) - (N-1)}, \quad \frac{(Q-1)(1-u)^2}{Q(1-u) - 1} \quad (97)$$

the second expression being the additional large- N approximation. At $(n, m) = (2, 1)$ the exact value is $12/17 \simeq 0.706$ against the approximation's 0.75. Both share the thermodynamic limit

$$\lim_{n \rightarrow \infty} \lim_{p \rightarrow 1^-} R_1(p) = 1 - u = 1 - 2^{-m}. \quad (98)$$




Figure 18: The exact one-question extraction ratio (95) of the spike prior, plotted only where it is defined ($1/N < p < 1$). Dashed lines are the exact finite- N limits (97) as $p \rightarrow 1^-$. Growing n drives them down to $1 - u = 1/2$; growing m lifts them toward one.

For the spike-plus-uniform family (98) is a theorem. Extending it to a broader class of uninformative hedges would require a precise definition and a separate extremal argument: full support alone is not sufficient.

A.3 The sharp-prior limit and the order of limits

The sharp-prior calculation takes $\varepsilon = 1 - p \rightarrow 0^+$ with n and m fixed. It is a limit toward the delta prior, not evaluation at $p = 1$, where leverage is $0/0$. With the entropy coefficient $\iota_k := \frac{N}{N-1}(1 - u^k)$ we have $P_k = 1 - \varepsilon \iota_k$, the other $u^{-k} - 1$ outcomes share mass $\varepsilon \iota_k$, and

$$G_k = h_2(\varepsilon \iota_k) + \varepsilon \iota_k \log_2(u^{-k} - 1) \sim \varepsilon \iota_k \log_2 \frac{1}{\varepsilon} \quad (\varepsilon \downarrow 0) \quad (99)$$

at fixed finite Q and k . Substituting into (51) cancels the common factor $\varepsilon \log_2(1/\varepsilon)$ and leaves, for $1 \leq k \leq Q$,

$$\lim_{p \rightarrow 1^-} \langle L_k \rangle = (1 - u) \frac{Q - (Q - k)u^k}{1 - u^k} = \underbrace{Q(1 - u)}_{\text{extensive plateau}} + \underbrace{(1 - u) \frac{k u^k}{1 - u^k}}_{O(1) \text{ boundary layer}}. \quad (100)$$

The second term decays geometrically in k and does not depend on Q at all. In particular $\lim_{p \rightarrow 1^-} \langle L_1 \rangle = Q - (Q - 1)2^{-m}$, so

$$\lim_{n \rightarrow \infty} \lim_{p \rightarrow 1^-} \frac{\langle L_1 \rangle}{Q} = 1 - 2^{-m}. \quad (101)$$

The coefficient $1 - 2^{-m}$ is the fraction of the hypothesis count eliminated by one answer. Its equality with the normalized leverage follows from the entropy asymptotics and should not be read as literally settling that fraction of every individual column. The limiting operation is necessary: at $p = 1$ the learner already assigns probability one to the truth, so received and destroyed information both vanish and leverage is $0/0$. The finite value is the ratio at which the two vanish as $p \rightarrow 1^-$. Taking $k = \lfloor tQ \rfloor$ in (100) gives the same constant at every $t > 0$, with the uniform bound

$$\sup_{1 \leq k \leq Q} \left| \frac{1}{Q} \lim_{p \rightarrow 1^-} \langle L_k \rangle - (1 - u) \right| = O(Q^{-1}), \quad (102)$$

so the microscopic boundary term survives only as an $O(1)$ correction to the unnormalized leverage.




Figure 19: Left: the culling coefficient (101) against the answer width. Middle: the sharp-prior curve (100) normalized by Q , whose $O(1)$ boundary layer is squeezed into the origin as n grows, leaving the plateau at $1 - u$. Right: the approach at fixed $n = 6$, $m = 1$, labelled by $1 - p$. Tightening the prior lifts the curve only logarithmically, so the plateau is approached pointwise in t but never uniformly.

As $p \rightarrow 1^-$ the fixed- p coefficient of (69) diverges, $[h_{\text{spike}}(p, m) - (1 - p)m] / [(1 - p)m] \sim \frac{1 - 2^{-m}}{m} \log_2 \frac{1}{1 - p}$, which is the noncommutation in explicit form. Writing $L_{p,n}(t) := \langle L_{\lfloor tQ \rfloor} \rangle$ and comparing under the same normalization,

$$\lim_{n \rightarrow \infty} \lim_{p \rightarrow 1^-} \frac{L_{p,n}(t)}{Q} = 1 - 2^{-m}, \quad \lim_{p \rightarrow 1^-} \lim_{n \rightarrow \infty} \frac{L_{p,n}(t)}{Q} = 0. \quad (103)$$

Before dividing by Q the first ordering is extensive while the second diverges only logarithmically in $1/(1 - p)$. The two orders are the endpoints of a family of joint scalings. Choose $p_n \rightarrow 1$ with prescribed exponential sharpness

$$\frac{\log_2(1/(1 - p_n))}{Q} \rightarrow \lambda \in (0, \infty), \quad (104)$$

so that $\lambda = 0$ recovers the scaling behind the second limit and the formal endpoint $\lambda = \infty$ the first. Putting $\varepsilon_n = 1 - p_n$, the leading entropies for $0 < t < 1$ and $k = \lfloor tQ \rfloor$ are $G_1 \sim (1 - u)\varepsilon_n \lambda Q$, $G_k \sim \varepsilon_n Q(\lambda + mx)$ and $G_{k+1} - G_k \sim \varepsilon_n m$, and substitution into (51) gives the crossover

$$\frac{L_{p_n,n}(t)}{Q} \rightarrow (1 - u) \frac{\lambda}{\lambda + mx}. \quad (105)$$

At $t = 1$ the increment $G_{k+1} - G_k$ is undefined and the same value follows instead from the completion formula. The two familiar orders are thus endpoints of a larger family.

B Block priors and the Bernstein profile language

A stable, projectively consistent exchangeable prior cannot produce a smooth non-hyperbolic macroscopic curve: de Finetti reduces it to an i.i.d. mixture, whose interior profile is constant and whose macroscopic curve is a hyperbola. To obtain a profile that changes throughout $0 < t < 1$, the prior must retain structure tied to particular groups or addresses of questions, and a fixed partition into local blocks is the simplest solvable construction. The statement

needs its consistency qualification: a separate finite exchangeable law can be chosen at every n without forming one infinite exchangeable family, and such a sequence need be neither an i.i.d. mixture nor possessed of any thermodynamic limit.

B.1 Independent blocks of questions

Partition $\mathcal{Q}$ into independent blocks drawn from a fixed finite set of types, writing $\mathcal{P}_n$ for the partition and indexing types by κ . At finite n let $K_{\kappa,n}$ count the type- κ blocks and set $w_{\kappa,n} := r_{\kappa}K_{\kappa,n}/Q$, the fraction of all questions belonging to them, assuming $w_{\kappa,n} \rightarrow w_{\kappa}$ with $\sum_{\kappa} w_{\kappa} = 1$. A type κ carries a size r_{κ} , a fixed joint distribution on its r_{κ} answers shared by every block of that type up to relabelling, and a mean subset-entropy profile defined as follows. Let $\mathcal{V} \subseteq \mathcal{Q}$ be one particular block of type κ , fix an ordering of it once and for all, and give every subset $\mathcal{J} \subseteq \mathcal{V}$ the inherited order, so that $\psi(\mathcal{J}) \in \mathcal{A}^{|\mathcal{J}|}$ is its random answer vector. For $j \in \{0, \dots, r_{\kappa}\}$ define

$$H_{\kappa}(j) := \binom{r_{\kappa}}{j}^{-1} \sum_{\substack{\mathcal{J} \subseteq \mathcal{V} \\ |\mathcal{J}|=j}} H(\psi(\mathcal{J})). \quad (106)$$

Here j is the *number of questions selected from the block*, not a question label, and $H(\psi(\mathcal{J}))$ measures the learner's uncertainty in their *joint answers*, not uncertainty about which questions were selected: the set $\mathcal{J}$ is known. No symmetry is required for this averaged definition, although every zoo member below has the stronger property of **entropy homogeneity**, meaning that $H(\psi(\mathcal{J}))$ itself depends only on $|\mathcal{J}|$. Full permutation invariance is sufficient for that but not necessary: MDS code blocks are entropy homogeneous even though their coordinate distribution is not invariant under every permutation.

The profile records how much answer uncertainty remains visible after looking at j locations inside a block, and its purpose is that successive differences have an operational meaning. Averaging the chain rule over a uniform i -subset $\mathcal{J}$ and then a fresh question drawn uniformly from $\mathcal{V} \setminus \mathcal{J}$ gives the within-block increments

$$\eta_{\kappa}(i) := H_{\kappa}(i+1) - H_{\kappa}(i) = \mathbb{E}_{\mathcal{J},q} [H(\psi(q) \mid \psi(\mathcal{J}))], \quad 0 \leq i \leq r_{\kappa} - 1, \quad (107)$$

the mean uncertainty in one fresh answer after i of its block partners have been answered. This is why the block-entropy profile is useful: it converts a joint law on an entire block into exactly the conditional entropies the learning curve needs. The endpoints make the meaning concrete, $H_{\kappa}(0) = 0$ while $H_{\kappa}(r_{\kappa}) = H(\psi(\mathcal{V}))$ is the entropy of the whole block taken jointly. Two independent fair binary answers have profile $(0, 1, 2)$ and increments $(1, 1)$; the same fair bit copied into both questions has profile $(0, 1, 1)$ and increments $(1, 0)$, since after either answer is known the other carries no remaining uncertainty.

For a fresh question in a block of size r_{κ} , the number of its $r_{\kappa} - 1$ partners already present in a uniformly random ℓ -set is hypergeometric, so exactly at finite Q , for $0 \leq \ell < Q$,

$$\gamma_{n,\ell} = \sum_{\kappa} w_{\kappa,n} \sum_{i=0}^{r_{\kappa}-1} \frac{\binom{r_{\kappa}-1}{i} \binom{Q-r_{\kappa}}{\ell-i}}{\binom{Q-1}{\ell}} \eta_{\kappa}(i). \quad (108)$$

Holding every r_{κ} fixed, taking $\ell = \lfloor tQ \rfloor$ and assuming $w_{\kappa,n} \rightarrow w_{\kappa}$, sampling without replacement approaches independent sampling,

$$\frac{\binom{r-1}{i} \binom{Q-r}{\ell-i}}{\binom{Q-1}{\ell}} \longrightarrow \underbrace{\binom{r-1}{i} t^i (1-t)^{r-1-i}}_{\beta_{r-1,i}(t)}, \quad (109)$$

which simultaneously defines the Bernstein basis polynomial $\beta_{r-1,i}$ and, probabilistically, is the limiting probability that exactly i of the fresh question's $r-1$ partners have been asked by macroscopic time t . The limiting profile is therefore


$$\gamma(t) = \sum_{\kappa} w_{\kappa} \sum_{i=0}^{r_{\kappa}-1} \beta_{r_{\kappa}-1,i}(t) \eta_{\kappa}(i). \quad (110)$$

At the start of the run no partner has been asked, so only the $i=0$ term survives and $h_0 = \gamma(0) = \sum_{\kappa} w_{\kappa} \eta_{\kappa}(0) = \sum_{\kappa} w_{\kappa} H_{\kappa}(1)$: the same block-entropy profile supplies both the initial marginal entropy and the entire macroscopic conditional-entropy curve.

To construct a sequence with prescribed fractions, choose integer block counts with $r_{\kappa} K_{\kappa,n}/Q \rightarrow w_{\kappa}$ and $\sum_{\kappa} r_{\kappa} K_{\kappa,n} \leq Q$, assigning only $o(Q)$ leftover sites, whose treatment does not affect the limit; at finite n any leftovers can be recorded as an additional block type, so (108) still includes every question and remains exact. If block sizes grow with n the proof changes, a useful sufficient condition for the hypergeometric-to-binomial approximation being $(\max_{\mathcal{V} \in \mathcal{P}_n} |\mathcal{V}|)^2/Q \rightarrow 0$. For countably many types, pointwise frequency convergence must be replaced by an ℓ^1 or tightness condition.

EXAMPLE

Two types tile the growing table: pairs ($r_1 = 2$, blue) and quads ($r_2 = 4$, orange), aiming at $w_1 = 3/8$ and $w_2 = 5/8$. Integer counts force detours. At $n = 1$ only a pair fits and at $n = 2$ only a quad, while at $Q = 8$ and $Q = 16$ the best leftover-free tilings still miss the target shares by $1/8$.



Each cell is one question and each rounded group is one block, so the counts $K_{\kappa,n}$ can be read off: $(K_1, K_2) = (1, 0), (0, 1), (2, 1), (2, 3), (6, 5)$. The pair share $w_{1,n} = 1, 0, \frac{1}{2}, \frac{1}{4}, \frac{3}{8}$ closes in on its target by halving steps, and at $n = 5$ the tiling is exact: $6 \times 2 + 5 \times 4 = 32 = Q$.

B.2 Block laws and their map priors

The profile machinery never displayed the prior itself. It is constructive: each block law prices the restriction of a map to its block, and independence across blocks multiplies the prices. A map ϕ_j restricts to the answer tuple $\phi_j(\mathcal{V})$ on block $\mathcal{V}$, and

$$p_j = \prod_{\mathcal{V} \in \mathcal{P}_n} \Pr(\psi(\mathcal{V}) = \phi_j(\mathcal{V})). \quad (111)$$

The classical block laws, with the factor each contributes to (111) and the curve it produces:

block law	a size- r block is sampled by	factor $\Pr(\psi(\mathcal{V}) = y)$	leverage
clique	one value $a_{\mathcal{V}} \sim \nu_{\mathcal{V}}$ copied to every question	$\nu_{\mathcal{V}}(a)$ if $y = (a, \dots, a)$, else 0	$L \equiv r$, exact at every n
parity (σ)	uniform on the coset $\bigoplus_q y_q = \sigma$	$2^{-(r-1)}$ on the coset, else 0	rises $1 \rightarrow r/(r-1)$, late
tilted parity	i.i.d. Bernoulli(θ) bits conditioned on the coset	$\theta^{ y }(1-\theta)^{r- y }/\zeta_{r,\sigma}$ on the coset, else 0	same shape, bent by $\varepsilon = 1 - 2\theta$
$[r, k]_A$ MDS	uniform on an MDS code ($A \geq r$)	A^{-k} on codewords, else 0	$L(1) = r/k$; interior peak when mixed with fresh sites
lapse (δ)	any law above, replaced by the uniform block with probability δ	$(1-\delta)P_{\mathcal{V}}(y) + \delta A^{-r}$	full support; $\gamma(1) > 0$

The supports multiply as well: cliques carry the $A^{Q/r}$ maps constant on every block, parity carries the $2^{(r-1)Q/r}$ maps landing in every coset, a code carries $A^{kQ/r}$ maps, and the lapse restores $p_j > 0$ everywhere, playing the role the top class played in Section 4.3.

EXAMPLE

Cliques at $(n, m) = (2, 1)$: partition the four questions by their high bit, $\mathcal{V}_1 = \{00, 01\}$ and $\mathcal{V}_2 = \{10, 11\}$, with $\nu_1 = (0.8, 0.2)$ and $\nu_2 = (0.55, 0.45)$ on the answers $(0, 1)$. A supported map must be constant on both blocks, so it can depend only on b_1 : of the sixteen maps of Table 4, only FALSE ($j = 0$), $\neg b_1$ ($j = 3$), b_1 ($j = 12$) and TRUE ($j = 15$) survive, with product weights $p_0 = 0.8 \times 0.55 = 0.44$, $p_3 = 0.11$, $p_{12} = 0.36$, $p_{15} = 0.09$, and $p_j = 0$ for the other twelve. One answer from either block settles that block's remaining question for free: $L \equiv 2$.

From any row of the table the route to leverage is the one already walked in Section 4.3. The block law fixes the subset entropies (106), their differences give the increments η (107), and the hypergeometric average (108) is exact at every finite n ; its binomial limit is the Bernstein profile (110). No capacity theorem is needed: blocks are bounded, so elementary hypergeometric-to-binomial convergence replaces it. The finite law (74) then gives $\langle L_\ell \rangle$ exactly at every n , and the master law (76) gives the curve, with $h_0 = \sum_{\kappa} w_{\kappa} \eta_{\kappa}(0)$.

B.3 Bernstein polynomials as a profile language

For a single block type, suppressing the type label, equation (110) reads

$$\gamma(t) = \eta(0)(1-t)^2 + 2\eta(1)t(1-t) + \eta(2)t^2 \quad (112)$$

at $r = 3$, and $\eta(0)(1-t)^3 + 3\eta(1)t(1-t)^2 + 3\eta(2)t^2(1-t) + \eta(3)t^3$ at $r = 4$. Some useful coefficient vectors, with h_{fresh} the entropy of one independent fresh answer and h_{clique} the entropy of the value copied through a clique:

block	$(\eta(0), \dots, \eta(r-1))$	profile
independent fresh answers	$(h_{\text{fresh}}, \dots, h_{\text{fresh}})$	h_{fresh}
clique	$(h_{\text{clique}}, 0, \dots, 0)$	$h_{\text{clique}}(1-t)^{r-1}$
parity	$(m, \dots, m, 0)$	$m(1-t^{r-1})$
(r, k) MDS	k entries m , then zeros	$m \Pr[\text{Bin}(r-1, t) \leq k-1]$

Entropy submodularity gives $m \geq \eta_\kappa(0) \geq \dots \geq \eta_\kappa(r_\kappa - 1) \geq 0$, and the derivative identity for Bernstein polynomials makes the macroscopic monotonicity explicit,

$$\gamma'(t) = \sum_{\kappa} w_{\kappa}(r_{\kappa} - 1) \sum_{i=0}^{r_{\kappa}-2} (\eta_{\kappa}(i+1) - \eta_{\kappa}(i)) \beta_{r_{\kappa}-2,i}(t) \leq 0, \quad (113)$$

so γ is nonincreasing as the general entropy calculus requires. Leverage itself need not be monotone.




Figure 20: The degree-three Bernstein basis: bump functions scanning left to right, the i -th peaked at $i/3$. A block profile is a positive combination of such bumps weighted by the microscopic conditional entropies $\eta(i)$, so the shape of L becomes a question about *where the coefficients drop*.

Equation (110) describes every prior in the stated class of independent blocks drawn from a fixed finite list of bounded types, and mixing types mixes their profiles linearly. There is also a useful approximation result when the answer alphabet is allowed to grow: mixtures of common-length MDS blocks realize any nonincreasing Bernstein coefficient sequence, and Bernstein polynomials approximate every continuous nonincreasing profile. A Reed–Solomon block of length r needs $A \geq r$, however, so sending $r \rightarrow \infty$ in that approximation also requires $m \rightarrow \infty$, or a different family of realizable entropy vectors. It is not a completeness theorem at fixed answer width.

B.4 Leverage curves at finite n

Because (108) is exact, the finite- n curves need no simulation: the hypergeometric increments feed the finite law (74) directly, and the limits below are the master law (76) applied to the Bernstein profiles of the table above. Three representative members, at $Q = 16$ and $Q = 64$ against their limits:




Figure 21: Exact finite- n leverage (dots, $Q = 16$ and 64) against the macroscopic limits (dashed), from (108) and (74) with no simulation. Left: cliques of four with a fair shared bit; the dots sit exactly on the constant $L \equiv 4$ at every finite n . Middle: parity blocks of four; deductions concentrate late and the finite curves approach the rise to $r/(r-1) = 4/3$ from below. Right: a 3:1 mixture of $[4, 2]$ MDS code blocks ($m = 2$) with fresh questions; the code releases its deductions around the threshold and the fresh sites keep receiving afterwards, an interior maximum of 1.645 at $t \simeq 0.72$ against the completion value 1.6.